

---

# Machine Learning Models and Predictive Analysis

Laura Lazarescou, John Rodgers, Maysam Mansor and Mel Schwan

March 1, 2021

---

## Introduction

All machine learning models are categorized as either supervised or unsupervised. Supervised learning involves learning a function that maps an input to an output based on example input-output pairs. Supervised models are then sub-categorized as a regression or classification model. In regression models, the result is continuous. Linear regression is merely finding a line that best fits the data. Decision trees are a popular model used in operations research, strategic planning, and machine learning. The last node of the decision tree, where a decision is made, is called the tree leaves. (Terence Shin - towards data science, 2020)

Extreme Gradient Boosting is a model that creates a partition tree to make predictions on class-level outcomes using data subsets. New, subsequent partition trees are applied to the remaining batches of the dataset until residual error is minimized. The weight of each sample batch is adaptively changed after each round of boosting (new tree). The model focuses on building trees to correctly explain data contributing to incorrect classifications. This is repeated until optimal performance is obtained. However, XGBoost is prone to over-fitting.

Support Vector Machine, is a supervised classification technique that can get pretty complicated but is intuitive at the most fundamental level. A support vector machine will find a hyperplane or a boundary between the the classes of data that maximizes the margin between the the classes. Many planes can separate the classes, but only one plane can maximize the margin or distance between the classes. This plane becomes the optimal solution for the model.

The third model type, Random Forest, is an ensemble learning technique that builds off of decision trees. Random forests involve creating multiple decision trees using bootstrapped datasets of the original data and randomly selecting a subset of variables at each decision tree step. Relying on a "majority wins" model reduces the risk of error from an individual tree.

We will train all three models in this study and vary the hyperparameters to maximize the models' accuracy values and minimize log-loss while considering the runtime requirements.

# Methods

## Data Preparation

Financial data, containing 538 features and a binary target are the corpus used for this study. There were no feature descriptions included. The features were labeled with non-descriptive letters like v1, v2 etc.


Our prepared dataset included 538 features once we applied one-hot encoding and categorical transformation to all non-numeric factors.

Factor v22 received exceptional treatment since it included over 18000 unique character values. Using a grouping approach, we categorized v22 into 53 factors based on the rationale that a value should have at least 125 occurrences to be represented in the dataset. If we had one-hot encoded v22 without this transformation, our final dataset would have included more than 19000 factors and it would have caused the feature space to grow intractably.

Sparse matrices can lead to poor models and results. Thus we have chosen to reduce the model feature set to the most impactful factors.

## Data Types and Data Distribution

The majority of our original factors were numeric. These factors are naturally compatible with machine learning models that we have used in this study. Figure 1 shows the unique values counts of non-numeric features, except for factor v22. It was removed from this plot because of its large number of unique values (18,210). v56 feature has the next greatest number of unique values.(125)



Unique Value Count per Non-Numeric Feature (excluding v22) (Figure 1)

## Model Validation

The models were validated using five(5)-fold cross-validation. The steps for model validation are given below. Hyperparameter tuning procedures and choices are discussed in each Model Discussion.

### Model Validation Procedure

- Split data into 67% training and 33% test sets. The train dataset consisted of 76595 rows and the test dataset included 37726 rows. Both datasets included 538 factors.
- Use the train dataset with cross-validation to build multiple models of each type.
- Estimate the log-loss of XGBoost and Random Forest models, and the accuracy for all models by using the test dataset as a validation dataset, then comparing the predicted target values to the known target values.
- Statistics and performance of each model are listed in Results.

# Models

In this case study, we use XGBoost, Support Vector Machines, and Random Forest to model the data. In each section we will describe how these models work, and what steps were taken to tune each model.

## Extreme Gradient Boosting

Extreme Gradient Boosting (XGBoost) is laser focused on computational speed and model performance. (Jason Brownlee; 2016, machine learning mastery)

### Model Features

Three main forms of gradient boosting are supported:

- Gradient Boosting algorithm also called gradient boosting machine including the learning rate.
- Stochastic Gradient Boosting with sub-sampling at the row, column and column per split levels.
- Regularized Gradient Boosting with both L1 and L2 regularization.

### Algorithm Features

Some key algorithm implementation features include:

- Sparse Aware implementation with automatic handling of missing data values.
- Block Structure to support the parallelization of tree construction.
- Continued Training so that you can further boost an already fitted model on new data.

## XGBoost Hyperparameter Tuning

To find an optimal combination of hyperparameters for an XGBoost model, a randomized search of combinations was performed to identify the best performing model based on the value of log loss. Each of these hyperparameter combinations was evaluated using 5-fold cross validation of the training data set. The following hyper-paramaters and values were incorporated into the randomized grid search. (Table 1)

Hyperparameter	Values
max_depth	6, 10, 15, 20
learning_rate	0.001, 0.01, 0.1, 0.2, 0.3
subsample	0.5, 0.6, 0.7, 0.8, 0.9, 1.0
colsample_bytree	0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0
colsample_bylevel	0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0
min_child_weight	0.5, 1.0, 3.0, 5.0, 7.0, 10.0
gamma	0, 0.25, 0.5, 1.0
reg_lambda	0.1, 1.0, 5.0, 10.0, 50.0, 100.0


**Table 1**

The search model selected 5 hyperparameter combinations at random from the list above. With each of these 5 models being evaluated with a 5 cross-fold cross-validation, a total of 25 models were evaluated to determine the best-performing combination of hyperparameters. Log loss was used to identify the best-performing model, with the following combination of hyperparameters returning a log-loss value of 0.4691 and an accuracy score of 0.7683. (Table 2)

Hyperparameter	Value
max_depth	10
learning_rate	0.1
subsample	0.9
colsample_bytree	0.5
colsample_bylevel	0.4
min_child_weight	7.0
gamma	0.5
reg_lambda	10.0

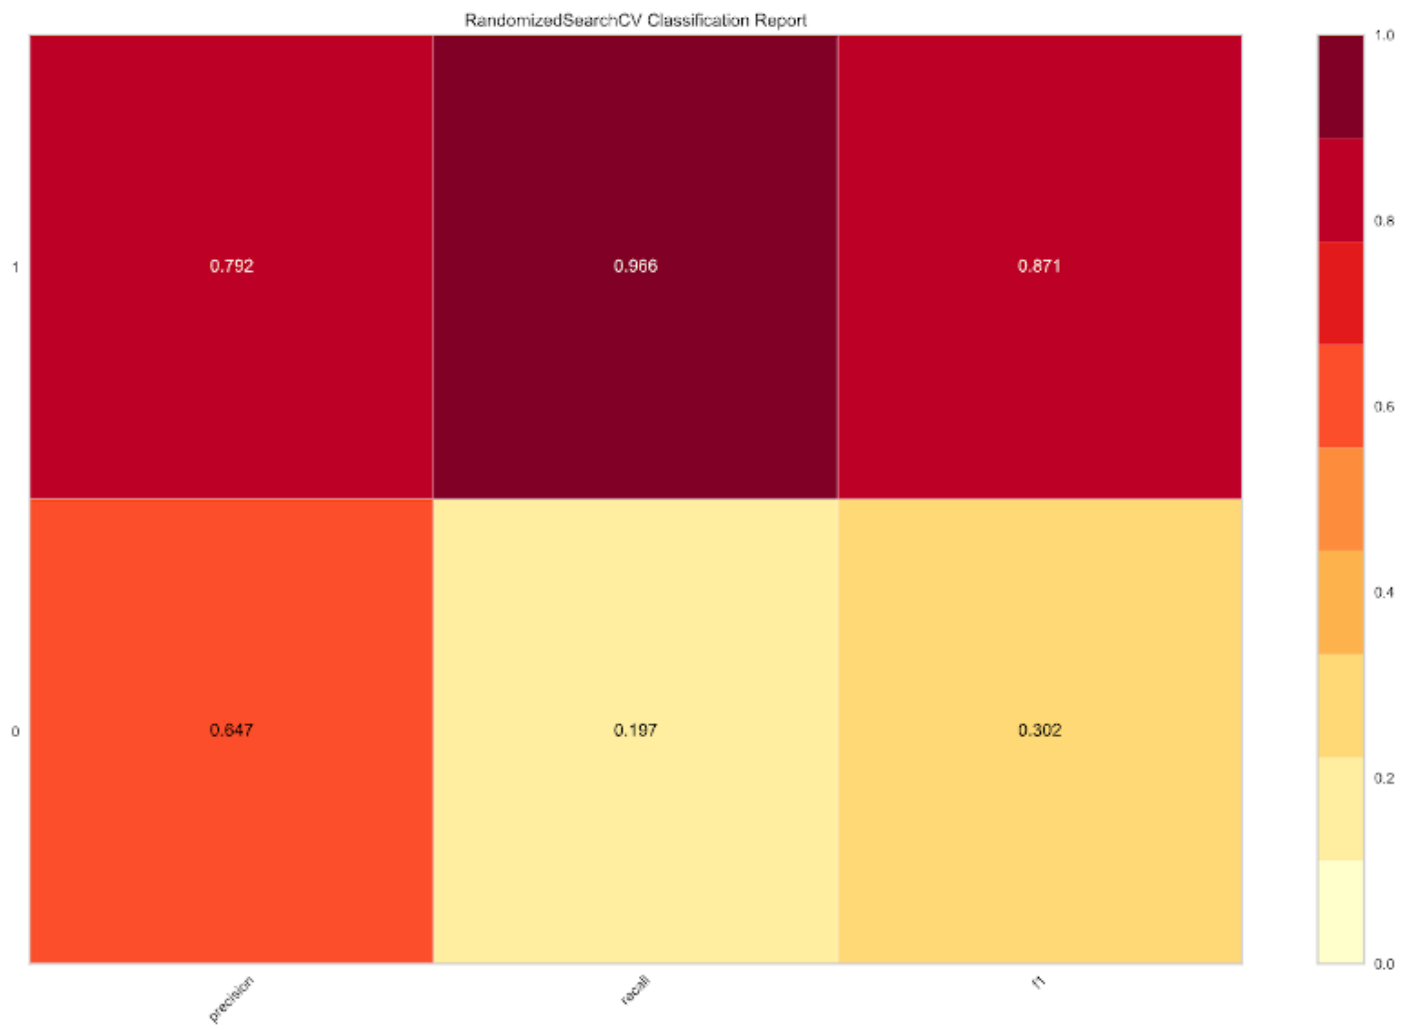
**Table 2**

Additional metrics for the tuned XGBoost model were also evaluated. The following figure shows the actual classes of the test data compared to the value that the model predicted. (Figure 2)



**XGB Class Prediction (Figure 2)**

The figure below shows the precision, recall, and f1 score for each of the two classes. (Figure 3)



**XGB Classification Report (Figure 3)**

## Support Vector Machine

A support vector machine is a supervised learning algorithm that sorts data into two categories. It is trained with a series of data already classified into two categories, building the model as it is initially trained. The task of an SVM algorithm is to determine which category a new data point belongs in. This makes SVM a kind of non-binary linear classifier.

An SVM algorithm should not only place objects into categories, but have the margins between them on a graph as wide as possible.

### Advantages:

- Works relatively well when there is a clear margin of separation between classes
- More effective in high dimensional spaces
- Effective in cases where the number of dimensions is greater than the number of samples
- Relatively memory efficient

### Disadvantages:

- Algorithm is not suitable for large data sets
- Does not perform very well when the data set has more noise i.e. target classes are overlapping
- In cases where the number of features for each data point exceeds the number of training data samples, the SVM will underperform
- Works by putting data points, above and below the classifying hyperplane there is no probabilistic explanation for the classification

Estimating the SVM in these high-dimensional spaces is considerably computationally expensive. Consider the model complexity when determining whether SVM should be implemented.

### SVM Hyperparameter Tuning

Hyperparameters were selected with a randomized search. `GridSearchCV` is a library function that is a member of sklearn's `model_selection` package. It loops through predefined Hyperparameter and fits the model on the training set. In the end, the best parameters from the list of hyperparameters can be selected. In our `LinearSVC` we evaluated hyperparameters (`C`, `loss`, `penalty`, `dual` and `tol`, `max_iter`) and set a range for each and at the end we got an accuracy of 0.77 with the best parameters selected and tabulated in the table which gained 0.001 for `C`, `squared_hinge` for `param_loss` and with `param_dual` become `False`, `tol = 1e-05` with 100 `max_iter` with 872 sec long and finally with an accuracy of 0.771137. We also hyper-tuned `LinearSVC` for sample sizes 1000, 2000, 5000, 10000 and found out the accuracy for each and the time of fitting the model was calculated and reported respectively. We showed results of parameter tuning and accuracy in corresponding tables.



## Random Forest

A Random Forest is an ensemble model created from a collection of decision trees and bootstrapped aggregated (bagged) data (Breiman, 1996; James et al, 2013). The following steps are used to create bagged trees:

- bootstrap sample (repeated sampling with replacement) the dataset to create  $B$  separate datasets.
- fit a model  $f^b(x)$  on each  $B$  dataset.

The bagged decision tree model is the majority vote of the classifiers resulting in the class prediction. Generally an ensemble should consist of a large number of decision trees. The number of decision trees was used as a hyperparameter and we tuned it with cross-validation.

### Random Forest Hyperparameter Tuning

For Random Forest we also used GridSearchCV to randomly select combinations of hyper parameters. Outside of the Grid, the value of `n_estimators` is one of the most influential parameters in Random Forest. In our base model we chose `n_estimators = 10` and ran it on the complete train dataset.

Below in figure 4 is a list of the default or base parameters used in our RF Base Model:

Parameters Used by Base Model:

```
{'bootstrap': True,
 'ccp_alpha': 0.0,
 'class_weight': None,
 'criterion': 'gini',
 'max_depth': None,
 'max_features': 'auto',
 'max_leaf_nodes': None,
 'max_samples': None,
 'min_impurity_decrease': 0.0,
 'min_impurity_split': None,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'n_estimators': 10,
 'n_jobs': None,
 'oob_score': False,
 'random_state': 123,
 'verbose': 0,
 'warm_start': False}
```

Random Forest Base Parameters (Figure 4)

## Grid Parameter Options

We did not include all possible parameters that could be tuned in our grid definition. We also limited `n_iterations` to 5000 and in our results, the algorithm only chose 1000. From this we learned that if one wants to test some extreme conditions or specific grid configurations, a random grid search may not be the best approach. Figure 4 shows the grid parameters that could be chosen at random by the `GridSearchCV` function (Figure 5).


Random Grid Parameters:

```
{'bootstrap': [True, False],  
 'max_depth': [10, 100, 200, 300, 400, 500, 600, 700, 800, 900, 1000],  
 'max_features': ['auto', 'sqrt'],  
 'min_samples_leaf': [1, 2, 4],  
 'min_samples_split': [2, 5, 10],  
 'n_estimators': [100, 500, 1000, 1500]}
```

**Random Forest Grid Parameters (Figure 5)**

## Top Features

We will not share the top features from all models that were run, however we did notice that the top ten features from different models were not consistent. The feature importance graph in Figure 4 is difficult to read because there are so many features. However, what it does show well is that no features has a significant percent of importance. The top 10 features out of 508 features represent approximately 30% of the influence in the model. This says that the variability of the model is high, and we have seen that in our different scenarios (Figure 6)



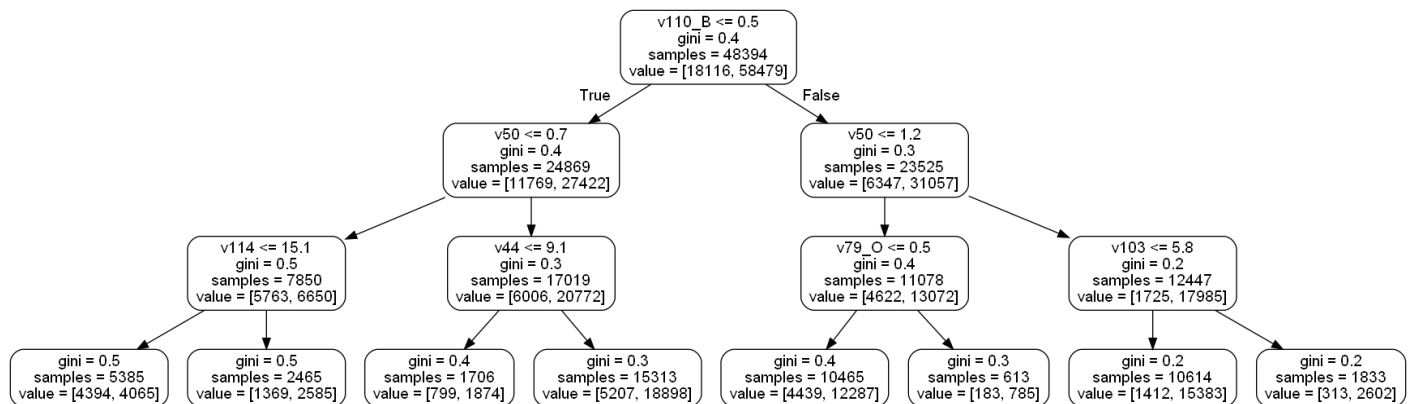
**Random Forest Feature Importance - All Features (Figure 6)**

## Top Ten Features

Rank	Feature ID	Importances
1.	feature 44	(0.060483)
2.	feature 11	(0.027440)
3.	feature 9	(0.025312)
4.	feature 96	(0.024937)
5.	feature 35	(0.023831)
6.	feature 20	(0.023625)
7.	feature 29	(0.023525)
8.	feature 13	(0.022473)
9.	feature 0	(0.021958)
10.	feature 338	(0.006959)

### \*\*Example Tree from our Random Forest Base Model"

The Figure below demonstrates a very small portion of the total Random Forest model. In this example, we intentionally limited the max\_depth of the tree to be able to visualize each element. (Figure 7)



Random Forest Tree with max\_depth=3 and n\_iterations=10 (Figure 7)

## Results

## Validation Results

Model results are provided in Table 4. Base models were constructed with the parameters that were provided by Dr. Slater. Hypertuned models are listed by type and variation. We were able to provide log-loss and accuracy for XGBoost and Random Forest. SVM does not provide log-loss so it is excluded from the table. In addition to log-loss and accuracy, we have provided timing or estimated timing for some models. This allows us to evaluate the cost/accuracy trade-offs.


Of the Random Forest models, our best performer was also the most simple and took the least amount of processing time because we did not use RandomizedSearchCV. The number of iterations for all randomized models was n\_iter=5 and the best\_params were all very similar between the different cases. It's very possible that the default parameters just happened to be the best performing combination of tunable parameters for this dataset. Given the relatively small amount of feature importance and the large number of features, it's also possible that the test dataset was very similar to the train dataset in its lack of correlation or internal trend. Whatever the reason, we find that random forest outperforms the XGBoost and SVM models with this data. (Table 3)

Model	Log-Loss	Accuracy	Wall Time (Seconds)
XGBoost (Pre-Tuned Model)	0.585068	0.768330	8
XGBoost (RandomizedSearchCV)	0.469189	0.781848	1534
SVM, Entire DSet	NA	0.7608	1179.79
SVM, For_1000	NA	0.779	13.2
SVM, For_2000	NA	0.4075	28.39
SVM, For_5000	NA	0.7662	71.83
SVM, For_10000	NA	0.7599	161.57
Random Forest Base	0.248184	0.927400	41
RF Tuned, 1000 Entries	0.483830	0.786000	120
RF Tuned, 5000 Entries	0.258795	0.917200	420
RF Tuned, Full Dataset	0.264958	0.913960	540

**Table of Model Performance and Results (Table 3)**

## AUC Comparisons for LinearSVC and Random Forest

Below in figure 1, is the comparison of the AUC performance for the LinearSVC and Random Forest model. This plot indicates that the Random Forest model trends towards more true positives than LinearSVC.decision\_function (Figure 8)



**LinearSVC and Random Forest Model Comparisons (Figure 8)**

## Conclusion

The best accuracy and log-loss values of our Random Forest (RF), Support Vector Machine (SVM), and XGBoost tuned models varied greatly depending on the hyperparameter tuning. Random Forest accuracy results were significantly higher than the other models. XGBoost and SVM were similar in their accuracy, but their execution time was significantly different. The main difference between all models appeared in the compute time required to ensure these results. The log-loss values for XGBoost were also higher than Random Forest. This difference in the predicted probability from the actual value in XGBoost and Random Forest makes our choice less difficult if you have the computing power. Random Forest appears to be the best model for this type of data set.

SVM is a useful model for small data sets that are highly dimensional. If you have a Big Data corpus, then XGBoost would be a good model. Random Forest works with categorical features very well and can handle high dimensional spaces and large numbers of training examples.

# References

Terence Shin (2020), towards data science - All Machine Learning Models Explained in 6 Minutes

Breiman, L. (1996). Bagging Predictors. Machine Learning, 24, 123-140.

CJason Brownlee (2016), machine learning mastery - A Gentle Introduction to XGBoost for Applied Machine Learning.

## Appendix

### Code

```
In [3]: import pandas as pd
import xgboost as xgb
import os
import time
import numpy as np
from sklearn.metrics import log_loss, accuracy_score
from sklearn.svm import SVC
from sklearn.svm import LinearSVC
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
from tabulate import tabulate
import pickle

from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import RandomizedSearchCV

from sklearn.tree import export_graphviz
from pprint import pprint

import math
import matplotlib.pyplot as plt
import matplotlib.pyplot as plt
from sklearn import model_selection

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.svm import LinearSVC
import sklearn.feature_selection as fs
from sklearn.model_selection import cross_val_score
```

```
In [7]: # Set Directory for image files - Comment out if you are not LL

import os
os.chdir('C:\\SMU_Local\\SMU_T5_QTW_CapA\\QTW_7333\\Unit 7 and 8\\case_s
tudy_81_2\\CS8')
```

```
In [8]: os.getcwd()
```

```
Out[8]: 'C:\\SMU_Local\\SMU_T5_QTW_CapA\\QTW_7333\\Unit 7 and 8\\case_study_81_
2\\CS8'
```

## Load Data and Prepare for Modeling

```
In [4]: # Load Data
# load data and separate target variable from dataset
train = pd.read_csv('Data/case_8.csv')
target = train['target']
train.drop(['target'],inplace=True, axis=1)
```

```
In [ ]: pickle.dump(target, open("Pickle/target.pkl", "wb"))
```

```
In [5]: # evaluate data types  
train.info(verbose=True)
```



```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 114321 entries, 0 to 114320
Data columns (total 132 columns):
#   Column  Dtype
---  -
0    ID      int64
1    v1      float64
2    v2      float64
3    v3      object
4    v4      float64
5    v5      float64
6    v6      float64
7    v7      float64
8    v8      float64
9    v9      float64
10   v10     float64
11   v11     float64
12   v12     float64
13   v13     float64
14   v14     float64
15   v15     float64
16   v16     float64
17   v17     float64
18   v18     float64
19   v19     float64
20   v20     float64
21   v21     float64
22   v22     object
23   v23     float64
24   v24     object
25   v25     float64
26   v26     float64
27   v27     float64
28   v28     float64
29   v29     float64
30   v30     object
31   v31     object
32   v32     float64
33   v33     float64
34   v34     float64
35   v35     float64
36   v36     float64
37   v37     float64
38   v38     int64
39   v39     float64
40   v40     float64
41   v41     float64
42   v42     float64
43   v43     float64
44   v44     float64
45   v45     float64
46   v46     float64
47   v47     object
48   v48     float64
49   v49     float64
50   v50     float64
51   v51     float64

```

52	v52	object
53	v53	float64
54	v54	float64
55	v55	float64
56	v56	object
57	v57	float64
58	v58	float64
59	v59	float64
60	v60	float64
61	v61	float64
62	v62	int64
63	v63	float64
64	v64	float64
65	v65	float64
66	v66	object
67	v67	float64
68	v68	float64
69	v69	float64
70	v70	float64
71	v71	object
72	v72	int64
73	v73	float64
74	v74	object
75	v75	object
76	v76	float64
77	v77	float64
78	v78	float64
79	v79	object
80	v80	float64
81	v81	float64
82	v82	float64
83	v83	float64
84	v84	float64
85	v85	float64
86	v86	float64
87	v87	float64
88	v88	float64
89	v89	float64
90	v90	float64
91	v91	object
92	v92	float64
93	v93	float64
94	v94	float64
95	v95	float64
96	v96	float64
97	v97	float64
98	v98	float64
99	v99	float64
100	v100	float64
101	v101	float64
102	v102	float64
103	v103	float64
104	v104	float64
105	v105	float64
106	v106	float64
107	v107	object
108	v108	float64

```

109 v109    float64
110 v110    object
111 v111    float64
112 v112    object
113 v113    object
114 v114    float64
115 v115    float64
116 v116    float64
117 v117    float64
118 v118    float64
119 v119    float64
120 v120    float64
121 v121    float64
122 v122    float64
123 v123    float64
124 v124    float64
125 v125    object
126 v126    float64
127 v127    float64
128 v128    float64
129 v129    int64
130 v130    float64
131 v131    float64
dtypes: float64(108), int64(5), object(19)
memory usage: 115.1+ MB

```

```

In [6]: # isolate object data type columns
train_object_dtype_cols = train.select_dtypes(include='object')

```

```

In [7]: # review head of object columns
train_object_dtype_cols.head()

```

```

Out[7]:


```

	v3	v22	v24	v30	v31	v47	v52	v56	v66	v71	v74	v75	v79	v91	v107	v110	v112	v1
0	C	XDX	C	C	A	C	G	DI	C	F	B	D	E	A	E	B	O	
1	C	GUV	C	C	A	E	G	DY	A	F	B	D	D	B	B	A	U	
2	C	FQ	E	C	A	C	F	AS	A	B	B	B	E	G	C	B	S	
3	C	ACUE	D	C	B	C	H	BW	A	F	B	D	B	B	B	B	J	
4	C	HIT	E	C	A	I	H	BW	C	F	B	D	C	G	C	A	T	

```
In [8]: # count unique values in object columns
train_object_dtype_cols.describe(include='all').loc['unique', :]
```

```
Out[8]: v3          3
v22     18210
v24          5
v30          7
v31          3
v47         10
v52         12
v56        122
v66          3
v71          9
v74          3
v75          4
v79         18
v91          7
v107         7
v110         3
v112         22
v113         36
v125        90
Name: unique, dtype: object
```

```
In [9]: # count unique values without v22 in object columns and plot
unique_series = train_object_dtype_cols.describe(include='all').loc['unique', :]
unique_df = pd.DataFrame(unique_series)
unique_df = unique_df.drop(['v22'])
ax = unique_df.plot.bar(title='Feature Unique Value Count')
```



In [52]: *# EDA on v22 feature*

```
v22_counts = train_object_dtype_cols.groupby('v22').v22.count()  
v22_counts.describe()  
v22_counts.median()  
v22_counts[v22_counts>125].describe()
```

Out[52]:

count	61.000000
mean	281.704918
std	431.560670
min	126.000000
25%	147.000000
50%	167.000000
75%	226.000000
max	2886.000000

Name: v22, dtype: float64

```
In [53]: # Get counts of unique values in v22
val = train_object_dtype_cols['v22'].value_counts()
# identify values with counts > 125
y = val[val < 125].index
# replace values with count < 125 with NaN
train_object_dtype_cols['v22'] = train_object_dtype_cols['v22'].replace
({x:math.nan for x in y})
# Output v22 information after removing values less than 125
train_object_dtype_cols.info()
train_object_dtype_cols.groupby('v22').v22.count().describe()
```

```
2021-02-28 22:04:04,974 [35816] WARNING py.warnings:110: [JupyterRequire] C:\Anaconda\lib\site-packages\ipykernel_launcher.py:6: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 114321 entries, 0 to 114320
Data columns (total 19 columns):
#   Column  Non-Null Count  Dtype
---  -
0    v3      114321 non-null   object
1    v22     17184 non-null    object
2    v24     114321 non-null   object
3    v30     114321 non-null   object
4    v31     114321 non-null   object
5    v47     114321 non-null   object
6    v52     114321 non-null   object
7    v56     114321 non-null   object
8    v66     114321 non-null   object
9    v71     114321 non-null   object
10   v74     114321 non-null   object
11   v75     114321 non-null   object
12   v79     114321 non-null   object
13   v91     114321 non-null   object
14   v107    114321 non-null   object
15   v110    114321 non-null   object
16   v112    114321 non-null   object
17   v113    114321 non-null   object
18   v125    114321 non-null   object
dtypes: object(19)
memory usage: 16.6+ MB
```

```
Out[53]: count      61.000000
mean      281.704918
std       431.560670
min       126.000000
25%       147.000000
50%       167.000000
75%       226.000000
max       2886.000000
Name: v22, dtype: float64
```

```
In [54]: # one-hot encode remaining object columns
object_one_hot_df = pd.get_dummies(data=train_object_dtype_cols)
# Output one-hot encodeing results
object_one_hot_df.info(verbose=True)
# get list of columns that were one-hot encoded
drop_cols = train_object_dtype_cols.columns
```



```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 114321 entries, 0 to 114320  
Data columns (total 425 columns):
```

#	Column	Dtype
0	v3_A	uint8
1	v3_B	uint8
2	v3_C	uint8
3	v22_AAPP	uint8
4	v22_ABF	uint8
5	v22_ABOF	uint8
6	v22_ACHJ	uint8
7	v22_ACWE	uint8
8	v22_ACXD	uint8
9	v22_ADDF	uint8
10	v22_ADGN	uint8
11	v22_ADMI	uint8
12	v22_ADMP	uint8
13	v22_AFOZ	uint8
14	v22_AFYU	uint8
15	v22_AGDF	uint8
16	v22_AGON	uint8
17	v22_AGZT	uint8
18	v22_AHE	uint8
19	v22_AJQ	uint8
20	v22_AMR	uint8
21	v22_AWT	uint8
22	v22_AXH	uint8
23	v22_BLE	uint8
24	v22_DJU	uint8
25	v22_EJC	uint8
26	v22_GBS	uint8
27	v22_GEB	uint8
28	v22_GEJ	uint8
29	v22_HDD	uint8
30	v22_HUU	uint8
31	v22_HZE	uint8
32	v22_JGY	uint8
33	v22_KLZ	uint8
34	v22_LIP	uint8
35	v22_MNZ	uint8
36	v22_MQE	uint8
37	v22_NGS	uint8
38	v22_NRT	uint8
39	v22_NWG	uint8
40	v22_NXE	uint8
41	v22_OFD	uint8
42	v22_PBC	uint8
43	v22_PFR	uint8
44	v22_PSE	uint8
45	v22_PTJ	uint8
46	v22_PTO	uint8
47	v22_PWR	uint8
48	v22_QKI	uint8
49	v22_QKP	uint8
50	v22_QVR	uint8
51	v22_RIC	uint8

52	v22_ROZ	uint8
53	v22_TG	uint8
54	v22_TVR	uint8
55	v22_UAG	uint8
56	v22_VVI	uint8
57	v22_VZF	uint8
58	v22_WFT	uint8
59	v22_WNI	uint8
60	v22_WRI	uint8
61	v22_YEP	uint8
62	v22_YGJ	uint8
63	v22_YOD	uint8
64	v24_A	uint8
65	v24_B	uint8
66	v24_C	uint8
67	v24_D	uint8
68	v24_E	uint8
69	v30_A	uint8
70	v30_B	uint8
71	v30_C	uint8
72	v30_D	uint8
73	v30_E	uint8
74	v30_F	uint8
75	v30_G	uint8
76	v31_A	uint8
77	v31_B	uint8
78	v31_C	uint8
79	v47_A	uint8
80	v47_B	uint8
81	v47_C	uint8
82	v47_D	uint8
83	v47_E	uint8
84	v47_F	uint8
85	v47_G	uint8
86	v47_H	uint8
87	v47_I	uint8
88	v47_J	uint8
89	v52_A	uint8
90	v52_B	uint8
91	v52_C	uint8
92	v52_D	uint8
93	v52_E	uint8
94	v52_F	uint8
95	v52_G	uint8
96	v52_H	uint8
97	v52_I	uint8
98	v52_J	uint8
99	v52_K	uint8
100	v52_L	uint8
101	v56_A	uint8
102	v56_AA	uint8
103	v56_AB	uint8
104	v56_AC	uint8
105	v56_AE	uint8
106	v56_AF	uint8
107	v56_AG	uint8
108	v56_AH	uint8

109	v56_AI	uint8
110	v56_AJ	uint8
111	v56_AK	uint8
112	v56_AL	uint8
113	v56_AM	uint8
114	v56_AN	uint8
115	v56_AO	uint8
116	v56_AP	uint8
117	v56_AR	uint8
118	v56_AS	uint8
119	v56_AT	uint8
120	v56_AU	uint8
121	v56_AV	uint8
122	v56_AW	uint8
123	v56_AX	uint8
124	v56_AY	uint8
125	v56_AZ	uint8
126	v56_B	uint8
127	v56_BA	uint8
128	v56_BC	uint8
129	v56_BD	uint8
130	v56_BE	uint8
131	v56_BF	uint8
132	v56_BG	uint8
133	v56_BH	uint8
134	v56_BI	uint8
135	v56_BJ	uint8
136	v56_BK	uint8
137	v56_BL	uint8
138	v56_BM	uint8
139	v56_BN	uint8
140	v56_BO	uint8
141	v56_BP	uint8
142	v56_BQ	uint8
143	v56_BR	uint8
144	v56_BS	uint8
145	v56_BT	uint8
146	v56_BU	uint8
147	v56_BV	uint8
148	v56_BW	uint8
149	v56_BX	uint8
150	v56_BY	uint8
151	v56_BZ	uint8
152	v56_C	uint8
153	v56_CA	uint8
154	v56_CB	uint8
155	v56_CC	uint8
156	v56_CD	uint8
157	v56_CE	uint8
158	v56_CF	uint8
159	v56_CG	uint8
160	v56_CH	uint8
161	v56_CI	uint8
162	v56_CJ	uint8
163	v56_CK	uint8
164	v56_CL	uint8
165	v56_CM	uint8

166	v56_CN	uint8
167	v56_CO	uint8
168	v56_CP	uint8
169	v56_CQ	uint8
170	v56_CS	uint8
171	v56_CT	uint8
172	v56_CV	uint8
173	v56_CW	uint8
174	v56_CX	uint8
175	v56_CY	uint8
176	v56_CZ	uint8
177	v56_D	uint8
178	v56_DA	uint8
179	v56_DB	uint8
180	v56_DC	uint8
181	v56_DD	uint8
182	v56_DE	uint8
183	v56_DF	uint8
184	v56_DG	uint8
185	v56_DH	uint8
186	v56_DI	uint8
187	v56_DJ	uint8
188	v56_DK	uint8
189	v56_DL	uint8
190	v56_DM	uint8
191	v56_DN	uint8
192	v56_DO	uint8
193	v56_DP	uint8
194	v56_DQ	uint8
195	v56_DR	uint8
196	v56_DS	uint8
197	v56_DT	uint8
198	v56_DU	uint8
199	v56_DV	uint8
200	v56_DW	uint8
201	v56_DX	uint8
202	v56_DY	uint8
203	v56_DZ	uint8
204	v56_E	uint8
205	v56_F	uint8
206	v56_G	uint8
207	v56_H	uint8
208	v56_I	uint8
209	v56_L	uint8
210	v56_M	uint8
211	v56_N	uint8
212	v56_O	uint8
213	v56_P	uint8
214	v56_Q	uint8
215	v56_R	uint8
216	v56_T	uint8
217	v56_U	uint8
218	v56_V	uint8
219	v56_W	uint8
220	v56_X	uint8
221	v56_Y	uint8
222	v56_Z	uint8

223	v66_A	uint8
224	v66_B	uint8
225	v66_C	uint8
226	v71_A	uint8
227	v71_B	uint8
228	v71_C	uint8
229	v71_D	uint8
230	v71_F	uint8
231	v71_G	uint8
232	v71_I	uint8
233	v71_K	uint8
234	v71_L	uint8
235	v74_A	uint8
236	v74_B	uint8
237	v74_C	uint8
238	v75_A	uint8
239	v75_B	uint8
240	v75_C	uint8
241	v75_D	uint8
242	v79_A	uint8
243	v79_B	uint8
244	v79_C	uint8
245	v79_D	uint8
246	v79_E	uint8
247	v79_F	uint8
248	v79_G	uint8
249	v79_H	uint8
250	v79_I	uint8
251	v79_J	uint8
252	v79_K	uint8
253	v79_L	uint8
254	v79_M	uint8
255	v79_N	uint8
256	v79_O	uint8
257	v79_P	uint8
258	v79_Q	uint8
259	v79_R	uint8
260	v91_A	uint8
261	v91_B	uint8
262	v91_C	uint8
263	v91_D	uint8
264	v91_E	uint8
265	v91_F	uint8
266	v91_G	uint8
267	v107_A	uint8
268	v107_B	uint8
269	v107_C	uint8
270	v107_D	uint8
271	v107_E	uint8
272	v107_F	uint8
273	v107_G	uint8
274	v110_A	uint8
275	v110_B	uint8
276	v110_C	uint8
277	v112_A	uint8
278	v112_B	uint8
279	v112_C	uint8

280	v112_D	uint8
281	v112_E	uint8
282	v112_F	uint8
283	v112_G	uint8
284	v112_H	uint8
285	v112_I	uint8
286	v112_J	uint8
287	v112_K	uint8
288	v112_L	uint8
289	v112_M	uint8
290	v112_N	uint8
291	v112_O	uint8
292	v112_P	uint8
293	v112_Q	uint8
294	v112_R	uint8
295	v112_S	uint8
296	v112_T	uint8
297	v112_U	uint8
298	v112_V	uint8
299	v113_A	uint8
300	v113_AA	uint8
301	v113_AB	uint8
302	v113_AC	uint8
303	v113_AD	uint8
304	v113_AE	uint8
305	v113_AF	uint8
306	v113_AG	uint8
307	v113_AH	uint8
308	v113_AI	uint8
309	v113_AJ	uint8
310	v113_AK	uint8
311	v113_B	uint8
312	v113_C	uint8
313	v113_D	uint8
314	v113_E	uint8
315	v113_F	uint8
316	v113_G	uint8
317	v113_H	uint8
318	v113_I	uint8
319	v113_J	uint8
320	v113_L	uint8
321	v113_M	uint8
322	v113_N	uint8
323	v113_O	uint8
324	v113_P	uint8
325	v113_Q	uint8
326	v113_R	uint8
327	v113_S	uint8
328	v113_T	uint8
329	v113_U	uint8
330	v113_V	uint8
331	v113_W	uint8
332	v113_X	uint8
333	v113_Y	uint8
334	v113_Z	uint8
335	v125_A	uint8
336	v125_AA	uint8

337	v125_AB	uint8
338	v125_AC	uint8
339	v125_AD	uint8
340	v125_AE	uint8
341	v125_AF	uint8
342	v125_AG	uint8
343	v125_AH	uint8
344	v125_AI	uint8
345	v125_AJ	uint8
346	v125_AK	uint8
347	v125_AL	uint8
348	v125_AM	uint8
349	v125_AN	uint8
350	v125_AO	uint8
351	v125_AP	uint8
352	v125_AQ	uint8
353	v125_AR	uint8
354	v125_AS	uint8
355	v125_AT	uint8
356	v125_AU	uint8
357	v125_AV	uint8
358	v125_AW	uint8
359	v125_AX	uint8
360	v125_AY	uint8
361	v125_AZ	uint8
362	v125_B	uint8
363	v125_BA	uint8
364	v125_BB	uint8
365	v125_BC	uint8
366	v125_BD	uint8
367	v125_BE	uint8
368	v125_BF	uint8
369	v125_BG	uint8
370	v125_BH	uint8
371	v125_BI	uint8
372	v125_BJ	uint8
373	v125_BK	uint8
374	v125_BL	uint8
375	v125_BM	uint8
376	v125_BN	uint8
377	v125_BO	uint8
378	v125_BP	uint8
379	v125_BQ	uint8
380	v125_BR	uint8
381	v125_BS	uint8
382	v125_BT	uint8
383	v125_BU	uint8
384	v125_BV	uint8
385	v125_BW	uint8
386	v125_BX	uint8
387	v125_BY	uint8
388	v125_BZ	uint8
389	v125_C	uint8
390	v125_CA	uint8
391	v125_CB	uint8
392	v125_CC	uint8
393	v125_CD	uint8

```
394 v125_CE      uint8
395 v125_CF      uint8
396 v125_CG      uint8
397 v125_CH      uint8
398 v125_CI      uint8
399 v125_CJ      uint8
400 v125_CK      uint8
401 v125_CL      uint8
402 v125_D       uint8
403 v125_E       uint8
404 v125_F       uint8
405 v125_G       uint8
406 v125_H       uint8
407 v125_I       uint8
408 v125_J       uint8
409 v125_K       uint8
410 v125_L       uint8
411 v125_M       uint8
412 v125_N       uint8
413 v125_O       uint8
414 v125_P       uint8
415 v125_Q       uint8
416 v125_R       uint8
417 v125_S       uint8
418 v125_T       uint8
419 v125_U       uint8
420 v125_V       uint8
421 v125_W       uint8
422 v125_X       uint8
423 v125_Y       uint8
424 v125_Z       uint8
dtypes: uint8(425)
memory usage: 46.3 MB
```

```
In [55]: # drop one-hot encoded columns from dataframe
train = train.drop(drop_cols, axis=1)
```



```
In [56]: # merge one-hot encoded columns to dataframe
frames = [train, object_one_hot_df]
train = pd.concat(frames,axis=1)
# Output train dataframe
train.info(verbose=True)
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 114321 entries, 0 to 114320  
Data columns (total 538 columns):
```

#	Column	Dtype
0	ID	int64
1	v1	float64
2	v2	float64
3	v4	float64
4	v5	float64
5	v6	float64
6	v7	float64
7	v8	float64
8	v9	float64
9	v10	float64
10	v11	float64
11	v12	float64
12	v13	float64
13	v14	float64
14	v15	float64
15	v16	float64
16	v17	float64
17	v18	float64
18	v19	float64
19	v20	float64
20	v21	float64
21	v23	float64
22	v25	float64
23	v26	float64
24	v27	float64
25	v28	float64
26	v29	float64
27	v32	float64
28	v33	float64
29	v34	float64
30	v35	float64
31	v36	float64
32	v37	float64
33	v38	int64
34	v39	float64
35	v40	float64
36	v41	float64
37	v42	float64
38	v43	float64
39	v44	float64
40	v45	float64
41	v46	float64
42	v48	float64
43	v49	float64
44	v50	float64
45	v51	float64
46	v53	float64
47	v54	float64
48	v55	float64
49	v57	float64
50	v58	float64
51	v59	float64

52	v60	float64
53	v61	float64
54	v62	int64
55	v63	float64
56	v64	float64
57	v65	float64
58	v67	float64
59	v68	float64
60	v69	float64
61	v70	float64
62	v72	int64
63	v73	float64
64	v76	float64
65	v77	float64
66	v78	float64
67	v80	float64
68	v81	float64
69	v82	float64
70	v83	float64
71	v84	float64
72	v85	float64
73	v86	float64
74	v87	float64
75	v88	float64
76	v89	float64
77	v90	float64
78	v92	float64
79	v93	float64
80	v94	float64
81	v95	float64
82	v96	float64
83	v97	float64
84	v98	float64
85	v99	float64
86	v100	float64
87	v101	float64
88	v102	float64
89	v103	float64
90	v104	float64
91	v105	float64
92	v106	float64
93	v108	float64
94	v109	float64
95	v111	float64
96	v114	float64
97	v115	float64
98	v116	float64
99	v117	float64
100	v118	float64
101	v119	float64
102	v120	float64
103	v121	float64
104	v122	float64
105	v123	float64
106	v124	float64
107	v126	float64
108	v127	float64

109	v128	float64
110	v129	int64
111	v130	float64
112	v131	float64
113	v3_A	uint8
114	v3_B	uint8
115	v3_C	uint8
116	v22_AAPP	uint8
117	v22_ABF	uint8
118	v22_ABOF	uint8
119	v22_ACHJ	uint8
120	v22_ACWE	uint8
121	v22_ACXD	uint8
122	v22_ADDF	uint8
123	v22_ADGN	uint8
124	v22_ADMI	uint8
125	v22_ADMP	uint8
126	v22_AFOZ	uint8
127	v22_AFYU	uint8
128	v22_AGDF	uint8
129	v22_AGON	uint8
130	v22_AGZT	uint8
131	v22_AHE	uint8
132	v22_AJQ	uint8
133	v22_AMR	uint8
134	v22_AWT	uint8
135	v22_AXH	uint8
136	v22_BLE	uint8
137	v22_DJU	uint8
138	v22_EJC	uint8
139	v22_GBS	uint8
140	v22_GEB	uint8
141	v22_GEJ	uint8
142	v22_HDD	uint8
143	v22_HUU	uint8
144	v22_HZE	uint8
145	v22_JGY	uint8
146	v22_KLZ	uint8
147	v22_LIP	uint8
148	v22_MNZ	uint8
149	v22_MQE	uint8
150	v22_NGS	uint8
151	v22_NRT	uint8
152	v22_NWG	uint8
153	v22_NXE	uint8
154	v22_OFD	uint8
155	v22_PBC	uint8
156	v22_PFR	uint8
157	v22_PSE	uint8
158	v22_PTJ	uint8
159	v22_PTO	uint8
160	v22_PWR	uint8
161	v22_QKI	uint8
162	v22_QKP	uint8
163	v22_QVR	uint8
164	v22_RIC	uint8
165	v22_ROZ	uint8

166	v22_TG	uint8
167	v22_TVR	uint8
168	v22_UAG	uint8
169	v22_VVI	uint8
170	v22_VZF	uint8
171	v22_WFT	uint8
172	v22_WNI	uint8
173	v22_WRI	uint8
174	v22_YEP	uint8
175	v22_YGJ	uint8
176	v22_YOD	uint8
177	v24_A	uint8
178	v24_B	uint8
179	v24_C	uint8
180	v24_D	uint8
181	v24_E	uint8
182	v30_A	uint8
183	v30_B	uint8
184	v30_C	uint8
185	v30_D	uint8
186	v30_E	uint8
187	v30_F	uint8
188	v30_G	uint8
189	v31_A	uint8
190	v31_B	uint8
191	v31_C	uint8
192	v47_A	uint8
193	v47_B	uint8
194	v47_C	uint8
195	v47_D	uint8
196	v47_E	uint8
197	v47_F	uint8
198	v47_G	uint8
199	v47_H	uint8
200	v47_I	uint8
201	v47_J	uint8
202	v52_A	uint8
203	v52_B	uint8
204	v52_C	uint8
205	v52_D	uint8
206	v52_E	uint8
207	v52_F	uint8
208	v52_G	uint8
209	v52_H	uint8
210	v52_I	uint8
211	v52_J	uint8
212	v52_K	uint8
213	v52_L	uint8
214	v56_A	uint8
215	v56_AA	uint8
216	v56_AB	uint8
217	v56_AC	uint8
218	v56_AE	uint8
219	v56_AF	uint8
220	v56_AG	uint8
221	v56_AH	uint8
222	v56_AI	uint8

223	v56_AJ	uint8
224	v56_AK	uint8
225	v56_AL	uint8
226	v56_AM	uint8
227	v56_AN	uint8
228	v56_AO	uint8
229	v56_AP	uint8
230	v56_AR	uint8
231	v56_AS	uint8
232	v56_AT	uint8
233	v56_AU	uint8
234	v56_AV	uint8
235	v56_AW	uint8
236	v56_AX	uint8
237	v56_AY	uint8
238	v56_AZ	uint8
239	v56_B	uint8
240	v56_BA	uint8
241	v56_BC	uint8
242	v56_BD	uint8
243	v56_BE	uint8
244	v56_BF	uint8
245	v56_BG	uint8
246	v56_BH	uint8
247	v56_BI	uint8
248	v56_BJ	uint8
249	v56_BK	uint8
250	v56_BL	uint8
251	v56_BM	uint8
252	v56_BN	uint8
253	v56_BO	uint8
254	v56_BP	uint8
255	v56_BQ	uint8
256	v56_BR	uint8
257	v56_BS	uint8
258	v56_BT	uint8
259	v56_BU	uint8
260	v56_BV	uint8
261	v56_BW	uint8
262	v56_BX	uint8
263	v56_BY	uint8
264	v56_BZ	uint8
265	v56_C	uint8
266	v56_CA	uint8
267	v56_CB	uint8
268	v56_CC	uint8
269	v56_CD	uint8
270	v56_CE	uint8
271	v56_CF	uint8
272	v56_CG	uint8
273	v56_CH	uint8
274	v56_CI	uint8
275	v56_CJ	uint8
276	v56_CK	uint8
277	v56_CL	uint8
278	v56_CM	uint8
279	v56_CN	uint8

280	v56_CO	uint8
281	v56_CP	uint8
282	v56_CQ	uint8
283	v56_CS	uint8
284	v56_CT	uint8
285	v56_CV	uint8
286	v56_CW	uint8
287	v56_CX	uint8
288	v56_CY	uint8
289	v56_CZ	uint8
290	v56_D	uint8
291	v56_DA	uint8
292	v56_DB	uint8
293	v56_DC	uint8
294	v56_DD	uint8
295	v56_DE	uint8
296	v56_DF	uint8
297	v56_DG	uint8
298	v56_DH	uint8
299	v56_DI	uint8
300	v56_DJ	uint8
301	v56_DK	uint8
302	v56_DL	uint8
303	v56_DM	uint8
304	v56_DN	uint8
305	v56_DO	uint8
306	v56_DP	uint8
307	v56_DQ	uint8
308	v56_DR	uint8
309	v56_DS	uint8
310	v56_DT	uint8
311	v56_DU	uint8
312	v56_DV	uint8
313	v56_DW	uint8
314	v56_DX	uint8
315	v56_DY	uint8
316	v56_DZ	uint8
317	v56_E	uint8
318	v56_F	uint8
319	v56_G	uint8
320	v56_H	uint8
321	v56_I	uint8
322	v56_L	uint8
323	v56_M	uint8
324	v56_N	uint8
325	v56_O	uint8
326	v56_P	uint8
327	v56_Q	uint8
328	v56_R	uint8
329	v56_T	uint8
330	v56_U	uint8
331	v56_V	uint8
332	v56_W	uint8
333	v56_X	uint8
334	v56_Y	uint8
335	v56_Z	uint8
336	v66_A	uint8

337	v66_B	uint8
338	v66_C	uint8
339	v71_A	uint8
340	v71_B	uint8
341	v71_C	uint8
342	v71_D	uint8
343	v71_F	uint8
344	v71_G	uint8
345	v71_I	uint8
346	v71_K	uint8
347	v71_L	uint8
348	v74_A	uint8
349	v74_B	uint8
350	v74_C	uint8
351	v75_A	uint8
352	v75_B	uint8
353	v75_C	uint8
354	v75_D	uint8
355	v79_A	uint8
356	v79_B	uint8
357	v79_C	uint8
358	v79_D	uint8
359	v79_E	uint8
360	v79_F	uint8
361	v79_G	uint8
362	v79_H	uint8
363	v79_I	uint8
364	v79_J	uint8
365	v79_K	uint8
366	v79_L	uint8
367	v79_M	uint8
368	v79_N	uint8
369	v79_O	uint8
370	v79_P	uint8
371	v79_Q	uint8
372	v79_R	uint8
373	v91_A	uint8
374	v91_B	uint8
375	v91_C	uint8
376	v91_D	uint8
377	v91_E	uint8
378	v91_F	uint8
379	v91_G	uint8
380	v107_A	uint8
381	v107_B	uint8
382	v107_C	uint8
383	v107_D	uint8
384	v107_E	uint8
385	v107_F	uint8
386	v107_G	uint8
387	v110_A	uint8
388	v110_B	uint8
389	v110_C	uint8
390	v112_A	uint8
391	v112_B	uint8
392	v112_C	uint8
393	v112_D	uint8



394	v112_E	uint8
395	v112_F	uint8
396	v112_G	uint8
397	v112_H	uint8
398	v112_I	uint8
399	v112_J	uint8
400	v112_K	uint8
401	v112_L	uint8
402	v112_M	uint8
403	v112_N	uint8
404	v112_O	uint8
405	v112_P	uint8
406	v112_Q	uint8
407	v112_R	uint8
408	v112_S	uint8
409	v112_T	uint8
410	v112_U	uint8
411	v112_V	uint8
412	v113_A	uint8
413	v113_AA	uint8
414	v113_AB	uint8
415	v113_AC	uint8
416	v113_AD	uint8
417	v113_AE	uint8
418	v113_AF	uint8
419	v113_AG	uint8
420	v113_AH	uint8
421	v113_AI	uint8
422	v113_AJ	uint8
423	v113_AK	uint8
424	v113_B	uint8
425	v113_C	uint8
426	v113_D	uint8
427	v113_E	uint8
428	v113_F	uint8
429	v113_G	uint8
430	v113_H	uint8
431	v113_I	uint8
432	v113_J	uint8
433	v113_L	uint8
434	v113_M	uint8
435	v113_N	uint8
436	v113_O	uint8
437	v113_P	uint8
438	v113_Q	uint8
439	v113_R	uint8
440	v113_S	uint8
441	v113_T	uint8
442	v113_U	uint8
443	v113_V	uint8
444	v113_W	uint8
445	v113_X	uint8
446	v113_Y	uint8
447	v113_Z	uint8
448	v125_A	uint8
449	v125_AA	uint8
450	v125_AB	uint8

451	v125_AC	uint8
452	v125_AD	uint8
453	v125_AE	uint8
454	v125_AF	uint8
455	v125_AG	uint8
456	v125_AH	uint8
457	v125_AI	uint8
458	v125_AJ	uint8
459	v125_AK	uint8
460	v125_AL	uint8
461	v125_AM	uint8
462	v125_AN	uint8
463	v125_AO	uint8
464	v125_AP	uint8
465	v125_AQ	uint8
466	v125_AR	uint8
467	v125_AS	uint8
468	v125_AT	uint8
469	v125_AU	uint8
470	v125_AV	uint8
471	v125_AW	uint8
472	v125_AX	uint8
473	v125_AY	uint8
474	v125_AZ	uint8
475	v125_B	uint8
476	v125_BA	uint8
477	v125_BB	uint8
478	v125_BC	uint8
479	v125_BD	uint8
480	v125_BE	uint8
481	v125_BF	uint8
482	v125_BG	uint8
483	v125_BH	uint8
484	v125_BI	uint8
485	v125_BJ	uint8
486	v125_BK	uint8
487	v125_BL	uint8
488	v125_BM	uint8
489	v125_BN	uint8
490	v125_BO	uint8
491	v125_BP	uint8
492	v125_BQ	uint8
493	v125_BR	uint8
494	v125_BS	uint8
495	v125_BT	uint8
496	v125_BU	uint8
497	v125_BV	uint8
498	v125_BW	uint8
499	v125_BX	uint8
500	v125_BY	uint8
501	v125_BZ	uint8
502	v125_C	uint8
503	v125_CA	uint8
504	v125_CB	uint8
505	v125_CC	uint8
506	v125_CD	uint8
507	v125_CE	uint8

```
508 v125_CF      uint8
509 v125_CG      uint8
510 v125_CH      uint8
511 v125_CI      uint8
512 v125_CJ      uint8
513 v125_CK      uint8
514 v125_CL      uint8
515 v125_D       uint8
516 v125_E       uint8
517 v125_F       uint8
518 v125_G       uint8
519 v125_H       uint8
520 v125_I       uint8
521 v125_J       uint8
522 v125_K       uint8
523 v125_L       uint8
524 v125_M       uint8
525 v125_N       uint8
526 v125_O       uint8
527 v125_P       uint8
528 v125_Q       uint8
529 v125_R       uint8
530 v125_S       uint8
531 v125_T       uint8
532 v125_U       uint8
533 v125_V       uint8
534 v125_W       uint8
535 v125_X       uint8
536 v125_Y       uint8
537 v125_Z       uint8
dtypes: float64(108), int64(5), uint8(425)
memory usage: 144.9 MB
```

```
In [19]: # Save training set to pickle
pickle.dump(train, open("Pickle/train.pkl", "wb"))
```

```
In [20]: # Load train pickle for consistent shared data across models
train = pickle.load( open("Pickle/train.pkl", "rb" ) )
```

```
In [22]: # create test/train split of data and target
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(train, target, test_
size=0.33, random_state=42)
```

```
In [24]: #### Save split train and test sets for consistant model testing and tim
e savings.
pickle.dump(X_train, open("Pickle/X_train.pkl", "wb"))
pickle.dump(y_train, open("Pickle/y_train.pkl", "wb"))
pickle.dump(X_test, open("Pickle/X_test.pkl", "wb"))
pickle.dump(y_test, open("Pickle/y_test.pkl", "wb"))
```

```

In [111]: ##### Load split train and test sets for consistant model testing and tim
           e savings.
X_train = pickle.load( open("Pickle/X_train.pkl", "rb"))
y_train = pickle.load( open("Pickle/y_train.pkl", "rb"))
X_test = pickle.load( open("Pickle/X_test.pkl", "rb"))
y_test = pickle.load( open("Pickle/y_test.pkl", "rb"))

In [ ]: #####
        # End of Data Preparation

        # Begin Models
        #####

In [ ]: #####
        # XGBoost Model
        #####

In [ ]: # BEGIN JR XGB RANDOMIZED SEARCH

In [ ]: # combine data into DMatrix for XGBoost
xgtrain = xgb.DMatrix(X_train.values, y_train.values)
xgtest = xgb.DMatrix(X_test.values, y_test.values)

In [ ]: clf = xgb.XGBClassifier()

In [ ]: param_grid = {
        'silent': [False],
        'max_depth': [6, 10, 15, 20],
        'learning_rate': [0.001, 0.01, 0.1, 0.2, 0.3],
        'subsample': [0.5, 0.6, 0.7, 0.8, 0.9, 1.0],
        'colsample_bytree': [0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0],
        'colsample_bylevel': [0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 1.0],
        'min_child_weight': [0.5, 1.0, 3.0, 5.0, 7.0, 10.0],
        'gamma': [0, 0.25, 0.5, 1.0],
        'reg_lambda': [0.1, 1.0, 5.0, 10.0, 50.0, 100.0],
        'n_estimators': [100]}

In [ ]: xgb_rs_clf = RandomizedSearchCV(clf, param_grid, n_iter=5,
                                         n_jobs=-1, verbose=2, cv=5,
                                         scoring='neg_log_loss', refit=True, random_s
                                         tate=123)

In [ ]: import time
        print("Randomized search..")
        search_time_start = time.time()
        xgb_rs_clf.fit(X_train, y_train)
        xgb_rs_elapsed = time.time() - search_time_start
        print("Randomized search time:", xgb_rs_elapsed)

In [ ]: pickle.dump(xgb_rs_elapsed, open("xgb_rs_elapsed.pkl", "wb"))

```

```
In [ ]: xgb_rs_elapsed = pickle.load( open("xgb_rs_elapsed.pkl", "rb" ) )
```

```
In [ ]: pickle.dump(xgb_rs_clf, open("xgb_rs_clf.pkl", "wb"))
```

```
In [ ]: xgb_rs_clf = pickle.load( open("xgb_rs_clf.pkl", "rb" ) )
```

```
In [ ]: print('The best combination of parameters, based on an evaluation of log  
-loss is:')  
xgb_rs_clf.best_params_
```

```
In [ ]: best_params = xgb_rs_clf.best_params_
```

```
In [ ]: print(f'The log-loss value for the best model is {round(xgb_rs_clf.best_  
score_ * -1,4)}.'.)
```

```
In [ ]: results = pd.DataFrame(xgb_rs_clf.cv_results_)
```

```
In [ ]: pickle.dump(results, open("results.pkl", "wb"))
```

```
In [ ]: results = pickle.load( open("results.pkl", "rb" ) )
```

```
In [ ]: results.iloc[xgb_rs_clf.best_index_]
```

```
In [ ]: xgb_test_preds = xgb_rs_clf.predict(X_test)
```

```
In [ ]: pickle.dump(xgb_test_preds, open("xgb_test_preds.pkl", "wb"))
```

```
In [ ]: xgb_test_preds = pickle.load( open("xgb_test_preds.pkl", "rb" ) )
```

```
In [ ]: print(accuracy_score(y_test,np rint(xgb_test_preds)))
```

```
In [ ]: results_summary = pd.DataFrame([['XGBoost', 'RandomSearchCV',xgb_rs_clf.  
best_score_ * -1,accuracy_score(y_test,np rint(xgb_test_preds)),xgb_rs_e  
lapsed ]],columns=['Model', 'Tuning', 'log loss', 'accuracy', 'time'])
```

```
In [ ]: ##### END JR XGB RANDOMIZED SEARCH
```

```
In [ ]: ##### BEGIN JR XGB ORIGINAL EXAMPLE
```

```
In [ ]: print('Fit the model...')
        # XGBoost params:
        xgboost_params = {
            "objective": "binary:logistic",
            "booster": "gbtree",
            "eval_metric": "logloss",
            "eta": 0.01,
            "subsample": 0.5,
            "colsample_bytree": 0.5,
            "max_depth": 3
        }
        boost_round = 50

        xgb_clf_start = time.time()
        xgb_clf = xgb.train(xgboost_params,xgtrain,num_boost_round=boost_round,verbose_eval=True,maximize=False)
        xgb_clf_elapsed = time.time() - xgb_clf_start
```

```
In [ ]: pickle.dump(xgb_clf_elapsed, open("xgb_clf_elapsed.pkl", "wb"))
```

```
In [ ]: xgb_clf_elapsed = pickle.load( open("xgb_clf_elapsed.pkl", "rb" ) )
```

```
In [ ]: #Make predict
        print('Predict...')

        xgb_test_preds_orig = clf.predict(xgtest, ntree_limit=clf.best_iteration
        )

        # Save results
```

```
In [ ]: pickle.dump(xgb_clf, open("xgb_clf.pkl", "wb"))
```

```
In [ ]: xgb_clf = pickle.load( open("xgb_clf.pkl", "rb" ) )
```

```
In [ ]: print(log_loss(y_test,xgb_test_preds_orig))
        print(accuracy_score(y_test,np rint(xgb_test_preds_orig)))
```

```
In [ ]: results_summary = results_summary.append(pd.DataFrame([[ 'XGBoost', 'Base'
        ,log_loss(y_test,xgb_test_preds_orig),accuracy_score(y_test,np rint(xgb_
        test_preds_orig)), xgb_clf_elapsed ]],columns=[ 'Model', 'Tuning', 'log l
        oss', 'accuracy', 'time'])))
```

```
In [ ]: results_summary = results_summary.reset_index()
```

```
In [ ]: results_summary
```

```
In [ ]: pickle.dump(results_summary, open("jr_results_summary.pkl", "wb"))
```

```
In [ ]: ##### END XGB XGB ORIGINAL EXAMPLE
```

```
In [ ]: ##### BEGIN XGB VISUALIZATIONS
```

```
In [ ]: import matplotlib.pyplot as plt
from xgboost import plot_tree
```

```
In [ ]: fig, ax = plt.subplots(figsize=(30, 30))
plot_tree(xgb_rs_clf.best_estimator_, num_trees=1, ax=ax)
plt.show()
```

```
In [ ]: from yellowbrick.classifier import ClassificationReport, ClassPrediction
Error
```

```
In [ ]: report = ClassificationReport(xgb_rs_clf, size=(1080, 720), classes=[0,1
])

report.score(X_test, y_test)
c = report.poof()
```

```
In [ ]: error = ClassPredictionError(xgb_rs_clf, size=(1080, 720), classes=[0,1
])

error.score(X_test, y_test)
e = error.poof()
```

```
In [ ]: ##### END XGB VISUALIZATIONS
```

```
In [ ]: #####
# SVM model
#####
```

```
In [146]: ##### Load split train and test sets for consistant model testing and tim
e savings.
X_train = pickle.load( open("Pickle/X_train.pkl", "rb"))
y_train = pickle.load( open("Pickle/y_train.pkl", "rb"))
X_test = pickle.load( open("Pickle/X_test.pkl", "rb"))
y_test = pickle.load( open("Pickle/y_test.pkl", "rb"))
```

```
In [142]: #####
          ###   Original SVM Code from Dr. Slater with Added Timing
          #####

          start = time.time()
          #####
          # WARNING THIS TAKES AN HOUR TO RUN #
          # Using LinearSVC for faster returns#
          #####

          svm = LinearSVC(verbose=True, random_state=42)
          svm.fit(X_train, y_train)

          end = time.time()
          LinearSVC_time=round((end-start),2)
          LinearSVC_time
```

```
[LibLinear]
```

```
2021-03-01 00:26:57,492 [35816] WARNING py.warnings:110: [JupyterRequire] C:\Anaconda\lib\site-packages\sklearn\svm\_base.py:977: Convergence
Warning: Liblinear failed to converge, increase the number of iterations.
      "the number of iterations.", ConvergenceWarning)
```

```
Out[142]: 55.44
```

```
In [144]: # Predict using the model and Measure Accuracy


          X_pred=svm.predict(X_test)

          #from sklearn.metrics import accuracy_score
          svm_base_accuracy = accuracy_score(y_test,X_pred)
          print(svm_base_accuracy)

          0.5046122037851879
```



```
In [145]: # Plot the SCM curve
svm_disp = plot_roc_curve(svm, X_test, y_test)
```



```
In [ ]: #####
##### Begin MM CODE
#####
```

```
In [ ]: # FULL DATASET TUNED SVM MODEL
# Ignore to save time - Picked Models are available

# Gridsearch to determine the value of C
param_grid = {'C': [0.001,0.01,0.1],
              'loss': ['hinge', 'squared_hinge'],
              'penalty': ['l2'],
              'dual': [True,False],
              'tol': [0.00001,0.01], #0.0001 is the Default
              'max_iter': [100],
              }

SVC_Linear = LinearSVC(random_state=42)
CV_svc = GridSearchCV(estimator = SVC_Linear, param_grid=param_grid, cv=
5, n_jobs=-1,verbose=1)
Start=time.time()
CV_svc_mod = CV_svc.fit(X_train, y_train)
Stop=time.time()
Time1=Stop-Start
Time1
pkl_filename = "Pickle/CV_SVM_Linear.pkl"
with open(pkl_filename, 'wb') as file:
    pickle.dump(CV_svc_mod, file)
```

```
In [ ]: Time1
```

```
In [112]: # Load CV_svc_mod
#with open("C://Users/18322/OneDrive - Southern Methodist University/Desktop/QOW/week7/case study 8/case_study_8l_2/CV_SVM_Linear.pkl", 'rb') as file:
#    CV_svc = pickle.load(file)

CV_svc_mod = pickle.load( open("Pickle/CV_SVM_Linear.pkl", "rb" ) )
```

```
In [113]: svc_gridsearch = pd.DataFrame(CV_svc_mod.cv_results_)
svc_columns = [
    "param_C",
    "param_loss",
    "param_dual",
    "param_tol",
    "param_max_iter",
    "mean_fit_time",
    "mean_test_score",
    "rank_test_score"
]
```

```
In [114]: svc_gridsearch[svc_columns].sort_values(by="rank_test_score").head(10)
```

Out[114]:

	param_C	param_loss	param_dual	param_tol	param_max_iter	mean_fit_time	mean_test_s
6	0.001	squared_hinge	False	1e-05	100	872.557506	0.77
2	0.001	squared_hinge	True	1e-05	100	94.016914	0.77
3	0.001	squared_hinge	True	0.01	100	77.806056	0.77
7	0.001	squared_hinge	False	0.01	100	119.676386	0.77
23	0.1	squared_hinge	False	0.01	100	145.308221	0.77
14	0.01	squared_hinge	False	1e-05	100	1386.344374	0.77
15	0.01	squared_hinge	False	0.01	100	269.671581	0.77
22	0.1	squared_hinge	False	1e-05	100	1003.456923	0.77
10	0.01	squared_hinge	True	1e-05	100	207.509046	0.77
11	0.01	squared_hinge	True	0.01	100	214.311796	0.77

```
In [ ]: print('Best Accuracy:', CV_svc_mod.best_score_)
```

```
In [115]: # Create Dataframe of 1000 Rows
Xtrain_1000 = pd.DataFrame.sample(X_train, n=1000, random_state=123)
ytrain_1000 = pd.DataFrame.sample(y_train,n=1000, random_state=123)
Xtest_1000 = pd.DataFrame.sample(X_test,n=1000, random_state=123)
ytest_1000 = pd.DataFrame.sample(y_test,n=1000, random_state=123)
print('Training Features 1000:', Xtrain_1000.shape)
print('Training Labels 1000:', ytrain_1000.shape)
print('Testing Features 1000:', Xtest_1000.shape)
print('Testing Labels 1000:', ytest_1000.shape)
```

```
Training Features 1000: (1000, 538)
Training Labels 1000: (1000,)
Testing Features 1000: (1000, 538)
Testing Labels 1000: (1000,)
```

```
In [137]: #1000 Rows
# Gridsearch to determine the value of C
param_grid = {'C': [0.001,0.01,0.1],
              'loss': ['hinge', 'squared_hinge'],
              'penalty' : ['l2'],
              'dual' : [True,False],
              'tol': [0.00001,0.01], #0.0001 is the Default
              'max_iter': [100],
              }

SVC_Linear = LinearSVC(random_state=42)
CV_svc = GridSearchCV(estimator = SVC_Linear, param_grid=param_grid, cv=
5, n_jobs=-1,verbose=1)
Start=time.time()
CV_svc_mod_1000 = CV_svc.fit(Xtrain_1000, ytrain_1000)
Stop=time.time()
Time2=Stop-Start
Time2
pkl_filename = "Pickle/CV_SVM_Linear_1000.pkl"
with open(pkl_filename, 'wb') as file:
    pickle.dump(CV_svc_mod_1000, file)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
```

```
[Parallel(n_jobs=-1)]: Done 52 tasks          | elapsed:    0.7s
```

```
[Parallel(n_jobs=-1)]: Done 120 out of 120    | elapsed:    1.3s finished
```

```
In [ ]: Time2
```

```
In [138]: #1000
#with open("C://Users/18322/OneDrive - Southern Methodist University/Des
ktop/QOW/week7/case study 8/case_study_81_2/CV_SVM_Linear_1000.pkl", 'r
b') as file:
#    CV_svc = pickle.load(file)

CV_svc_mod_1000 = pickle.load( open("Pickle/CV_SVM_Linear_1000.pkl", "r
b" ) )
```

```
In [139]: svc_gridsearch = pd.DataFrame(CV_svc_mod_1000.cv_results_)
svc_columns = [
    "param_C",
    "param_loss",
    "param_dual",
    "param_tol",
    "param_max_iter",
    "mean_fit_time",
    "mean_test_score",
    "rank_test_score"
]
```

```
In [140]: #for 1000 samples
svc_gridsearch[svc_columns].sort_values(by="rank_test_score").head(10)
```

Out[140]:

	param_C	param_loss	param_dual	param_tol	param_max_iter	mean_fit_time	mean_test_s
23	0.1	squared_hinge	False	0.01	100	0.044880	0
15	0.01	squared_hinge	False	0.01	100	0.087365	0
7	0.001	squared_hinge	False	0.01	100	0.058641	0
6	0.001	squared_hinge	False	1e-05	100	0.117089	0
14	0.01	squared_hinge	False	1e-05	100	0.164079	0
22	0.1	squared_hinge	False	1e-05	100	0.122875	0
10	0.01	squared_hinge	True	1e-05	100	0.158082	0
11	0.01	squared_hinge	True	0.01	100	0.165455	0
1	0.001	hinge	True	0.01	100	0.156181	0
0	0.001	hinge	True	1e-05	100	0.143215	0

```
In [ ]: print('Best Accuracy:', CV_svc_mod_1000.best_score_)
```

```
In [121]: #2000
# Create Dataframe of 2000 Rows
Xtrain_2000 = pd.DataFrame.sample(X_train, n=2000, random_state=123)
ytrain_2000 = pd.DataFrame.sample(y_train, n=2000, random_state=123)
Xtest_2000 = pd.DataFrame.sample(X_test, n=2000, random_state=123)
ytest_2000 = pd.DataFrame.sample(y_test, n=2000, random_state=123)
print('Training Features 2000:', Xtrain_2000.shape)
print('Training Labels 2000:', ytrain_2000.shape)
print('Testing Features 2000:', Xtest_2000.shape)
print('Testing Labels 2000:', ytest_2000.shape)
```

```
Training Features 2000: (2000, 538)
Training Labels 2000: (2000,)
Testing Features 2000: (2000, 538)
Testing Labels 2000: (2000,)
```

```

In [122]: #2000
# Gridsearch to determine the value of C
param_grid = {'C': [0.001,0.01,0.1],
              'loss': ['hinge', 'squared_hinge'],
              'penalty' : ['l2'],
              'dual' : [True,False],
              'tol': [0.00001,0.01], #0.0001 is the Default
              'max_iter': [100],
              }

SVC_Linear = LinearSVC(random_state=42)
CV_svc = GridSearchCV(estimator = SVC_Linear, param_grid=param_grid, cv=
5, n_jobs=-1,verbose=1)
Start=time.time()
CV_svc_mod_2000 = CV_svc.fit(Xtrain_2000, ytrain_2000)
Stop=time.time()
Time3=Stop-Start
Time3
pkl_filename = "Pickle/CV_SVM_Linear_2000.pkl"
with open(pkl_filename, 'wb') as file:
    pickle.dump(CV_svc_mod_2000, file)

```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent work
ers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed:    8.5s
[Parallel(n_jobs=-1)]: Done 120 out of 120 | elapsed:   10.6s finished
2021-03-01 00:17:45,154 [35816] WARNING py.warnings:110: [JupyterRequi
re] C:\Anaconda\lib\site-packages\sklearn\svm\_base.py:977: Convergence
Warning: Liblinear failed to converge, increase the number of iteration
s.
    "the number of iterations.", ConvergenceWarning)

```

```

In [ ]: Time3

```

```

In [123]: #2000
#with open("C://Users/18322/OneDrive - Southern Methodist University/Des
ktop/QOW/week7/case study 8/case_study_81_2/CV_SVM_Linear_2000.pkl", 'r
b') as file:
#    CV_svc = pickle.load(file)

CV_svc_mod_2000 = pickle.load( open("Pickle/CV_SVM_Linear_2000.pkl", "r
b" ) )

```

```
In [124]: svc_gridsearch = pd.DataFrame(CV_svc_mod_2000.cv_results_)
svc_columns = [
    "param_C",
    "param_loss",
    "param_dual",
    "param_tol",
    "param_max_iter",
    "mean_fit_time",
    "mean_test_score",
    "rank_test_score"
]
```

```
In [126]: svc_gridsearch[svc_columns].sort_values(by="rank_test_score").head(10)
```

Out[126]:

	param_C	param_loss	param_dual	param_tol	param_max_iter	mean_fit_time	mean_test_s
6	0.001	squared_hinge	False	1e-05	100	0.417046	0.
23	0.1	squared_hinge	False	0.01	100	0.132361	0.
15	0.01	squared_hinge	False	0.01	100	0.134945	0.
7	0.001	squared_hinge	False	0.01	100	0.139359	0.
14	0.01	squared_hinge	False	1e-05	100	0.567534	0.
22	0.1	squared_hinge	False	1e-05	100	0.492509	0.
19	0.1	squared_hinge	True	0.01	100	0.442345	0.
18	0.1	squared_hinge	True	1e-05	100	0.464916	0.
17	0.1	hinge	True	0.01	100	0.449724	0.
16	0.1	hinge	True	1e-05	100	0.507707	0.

```
In [ ]: print('Best Accuracy:', CV_svc_mod_2000.best_score_)
```

```
In [127]: #5000
# Create Dataframe of 5000 Rows
Xtrain_5000 = pd.DataFrame.sample(X_train, n=5000, random_state=123)
ytrain_5000 = pd.DataFrame.sample(y_train, n=5000, random_state=123)
Xtest_5000 = pd.DataFrame.sample(X_test, n=5000, random_state=123)
ytest_5000 = pd.DataFrame.sample(y_test, n=5000, random_state=123)
print('Training Features 5000:', Xtrain_5000.shape)
print('Training Labels 5000:', ytrain_5000.shape)
print('Testing Features 5000:', Xtest_5000.shape)
print('Testing Labels 5000:', ytest_5000.shape)
```

```
Training Features 5000: (5000, 538)
Training Labels 5000: (5000,)
Testing Features 5000: (5000, 538)
Testing Labels 5000: (5000,)
```

```

In [ ]: # Gridsearch to determine the value of C
param_grid = {'C': [0.001,0.01,0.1],
              'loss': ['hinge', 'squared_hinge'],
              'penalty' : ['l2'],
              'dual' : [True,False],
              'tol': [0.00001,0.01], #0.0001 is the Default
              'max_iter': [100],
              }

SVC_Linear = LinearSVC(random_state=42)
CV_svc = GridSearchCV(estimator = SVC_Linear, param_grid=param_grid, cv=
5, n_jobs =-1,verbose=1)
Start=time.time()
CV_svc_mod_5000 = CV_svc.fit(Xtrain_5000, ytrain_5000)
Stop=time.time()
Time4=Stop-Start
Time4
pkl_filename = "Pickle/CV_SVM_Linear_5000.pkl"
with open(pkl_filename, 'wb') as file:
    pickle.dump(CV_svc_mod_5000, file)

```

```

In [ ]: Time4

```

```

In [128]: #with open("C://Users/18322/OneDrive - Southern Methodist University/Des
ktop/QOW/week7/case study 8/case_study_81_2/CV_SVM_Linear_5000.pkl", 'r
b') as file:
#     CV_svc = pickle.load(file)

CV_svc_mod_5000 = pickle.load( open("Pickle/CV_SVM_Linear_5000.pkl", "r
b" ) )

```

```

In [129]: svc_gridsearch = pd.DataFrame(CV_svc_mod_5000.cv_results_)
svc_columns = [
    "param_C",
    "param_loss",
    "param_dual",
    "param_tol",
    "param_max_iter",
    "mean_fit_time",
    "mean_test_score",
    "rank_test_score"
]

```

```
In [130]: svc_gridsearch[svc_columns].sort_values(by="rank_test_score").head(10)
```

Out[130]:

	param_C	param_loss	param_dual	param_tol	param_max_iter	mean_fit_time	mean_test_s
6	0.001	squared_hinge	False	1e-05	100	1.029846	0.0
22	0.1	squared_hinge	False	1e-05	100	0.939270	0.0
23	0.1	squared_hinge	False	0.01	100	0.230186	0.0
15	0.01	squared_hinge	False	0.01	100	0.244742	0.0
7	0.001	squared_hinge	False	0.01	100	0.227392	0.0
14	0.01	squared_hinge	False	1e-05	100	1.177088	0.0
0	0.001	hinge	True	1e-05	100	1.197936	0.0
1	0.001	hinge	True	0.01	100	1.296276	0.0
8	0.01	hinge	True	1e-05	100	1.226919	0.0
9	0.01	hinge	True	0.01	100	1.168087	0.0

```
In [ ]: print('Best Accuracy:', CV_svc_mod_5000.best_score_)
```

```
In [131]: #10000
# Create Dataframe of 10000 Rows
Xtrain_10000 = pd.DataFrame.sample(X_train, n=10000, random_state=123)
ytrain_10000 = pd.DataFrame.sample(y_train,n=10000, random_state=123)
Xtest_10000 = pd.DataFrame.sample(X_test,n=10000, random_state=123)
ytest_10000 = pd.DataFrame.sample(y_test,n=10000, random_state=123)
print('Training Features 10000:', Xtrain_10000.shape)
print('Training Labels 10000:', ytrain_10000.shape)
print('Testing Features 10000:', Xtest_10000.shape)
print('Testing Labels 10000:', ytest_10000.shape)

Training Features 10000: (10000, 538)
Training Labels 10000: (10000,)
Testing Features 10000: (10000, 538)
Testing Labels 10000: (10000,)
```



```
In [133]: # Gridsearch to determine the value of C
param_grid = {'C': [0.001,0.01,0.1],
              'loss': ['hinge', 'squared_hinge'],
              'penalty' : ['l2'],
              'dual' : [True,False],
              'tol': [0.00001,0.01], #0.0001 is the Default
              'max_iter': [100],
              }

SVC_Linear = LinearSVC(random_state=42)
CV_svc = GridSearchCV(estimator = SVC_Linear, param_grid=param_grid, cv=
5, n_jobs=-1,verbose=1)
Start=time.time()
CV_svc_mod_10000 = CV_svc.fit(Xtrain_10000, ytrain_10000)
Stop=time.time()
Time5=Stop-Start
Time5
pkl_filename = "Pickle/CV_SVM_Linear_10000.pkl"
with open(pkl_filename, 'wb') as file:
    pickle.dump(CV_svc_mod_10000, file)
```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent work  
ers.

[Parallel(n\_jobs=-1)]: Done 34 tasks | elapsed: 4.5s

[Parallel(n\_jobs=-1)]: Done 120 out of 120 | elapsed: 14.0s finished

```
In [ ]: Time5
```

```
In [134]: #with open("C://Users/18322/OneDrive - Southern Methodist University/Des
ktop/QOW/week7/case study 8/case_study_81_2/CV_SVM_Linear_10000.pkl", 'r
b') as file:
#     CV_svc = pickle.load(file)

CV_svc_mod_10000 = pickle.load( open("Pickle/CV_SVM_Linear_10000.pkl",
"rb" ) )
```

```
In [135]: svc_gridsearch = pd.DataFrame(CV_svc_mod_10000.cv_results_)
svc_columns = [
    "param_C",
    "param_loss",
    "param_dual",
    "param_tol",
    "param_max_iter",
    "mean_fit_time",
    "mean_test_score",
    "rank_test_score"
]
```

```
In [136]: svc_gridsearch[svc_columns].sort_values(by="rank_test_score").head(10)
```

Out[136]:

	param_C	param_loss	param_dual	param_tol	param_max_iter	mean_fit_time	mean_test_s
14	0.01	squared_hinge	False	1e-05	100	1.248989	0.
22	0.1	squared_hinge	False	1e-05	100	1.117655	0.
6	0.001	squared_hinge	False	1e-05	100	1.184583	0.
23	0.1	squared_hinge	False	0.01	100	0.249406	0.
15	0.01	squared_hinge	False	0.01	100	0.241507	0.
7	0.001	squared_hinge	False	0.01	100	0.262091	0.
19	0.1	squared_hinge	True	0.01	100	1.260491	0.
18	0.1	squared_hinge	True	1e-05	100	1.278927	0.
10	0.01	squared_hinge	True	1e-05	100	1.309391	0.
11	0.01	squared_hinge	True	0.01	100	1.259128	0.

```
In [ ]: print('Best Accuracy:', CV_svc_mod_10000.best_score_)
```

```
In [5]: #####  
# Random Forest Model  
#####
```

```

In [6]: #####
# RF Base Model
#####

#from sklearn.ensemble import RandomForestClassifier

# Full Dataset used for Base Model.
# Default max_levels is None, so the tree is very large

rf_base = RandomForestClassifier(n_estimators=10, random_state=123 )

start = time.time()
rf_base.fit(X_train, y_train)

end = time.time()
rf_base_time=round((end-start),2)
rf_base_time

-----
----
NameError                                Traceback (most recent call l
ast)
<ipython-input-6-cda592a95acd> in <module>
      8 # Default max_levels is None, so the tree is very large
      9
--> 10 rf_base = RandomForestClassifier(n_estimators=10, random_state=
123 )
     11
     12 start = time.time()

NameError: name 'RandomForestClassifier' is not defined

```

```

In [63]: rf_base_preds = rf_base.predict_proba(X_test)

```

```

In [64]: rf_base_log_loss = log_loss(y_test,rf_base_preds[:,1]) # each column is
class probability,
print(rf_base_log_loss)
rf_base_accuracy = accuracy_score(y_test,np.rint(rf_base_preds[:,1]))
print(rf_base_accuracy)

0.9337292313310098
0.7515241478025765

```

```

In [65]: #####
# Begin Hypertuning for Random Forest
#####

```

```
In [66]: # Citation: Thank you to Will Koehrsen and Towards Data Science. RF Hypertuning code is borrowed from his example.
# https://towardsdatascience.com/hyperparameter-tuning-the-random-forest-in-python-using-scikit-learn-28d2aa77dd74

# Note - different from tutorial, we are using rf=RandomForestClassifier(), not RandomForestRegressor()
from pprint import pprint
# Look at parameters used by our current forest
print('Parameters Used by Base Model:\n')
pprint(rf_base.get_params())
```

Parameters Used by Base Model:

```
{'bootstrap': True,
 'ccp_alpha': 0.0,
 'class_weight': None,
 'criterion': 'gini',
 'max_depth': None,
 'max_features': 'auto',
 'max_leaf_nodes': None,
 'max_samples': None,
 'min_impurity_decrease': 0.0,
 'min_impurity_split': None,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'n_estimators': 10,
 'n_jobs': None,
 'oob_score': False,
 'random_state': 123,
 'verbose': 0,
 'warm_start': False}
```

```
In [67]: # Define Random Grid Parameters and use RandomizedSearchCV to choose
# different combinations of parameters for different sizes of datasets

#from sklearn.model_selection import RandomizedSearchCV
# Number of trees in random forest
n_estimators = [100,500,1000,1500]
# Number of features to consider at every split
max_features = ['auto', 'sqrt']
# Maximum number of levels in tree
max_depth = [10, 100, 200,300,400,500,600,700,800,900,1000]
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Method of selecting samples for training each tree
bootstrap = [True, False]
# Create the random grid
random_grid = {'n_estimators': n_estimators,
               'max_features': max_features,
               'max_depth': max_depth,
               'min_samples_split': min_samples_split,
               'min_samples_leaf': min_samples_leaf,
               'bootstrap': bootstrap}
print('Random Grid Parameters:\n')
pprint(random_grid)
```

Random Grid Parameters:

```
{'bootstrap': [True, False],
 'max_depth': [10, 100, 200, 300, 400, 500, 600, 700, 800, 900, 1000],
 'max_features': ['auto', 'sqrt'],
 'min_samples_leaf': [1, 2, 4],
 'min_samples_split': [2, 5, 10],
 'n_estimators': [100, 500, 1000, 1500]}
```

```
In [68]: # Adapt variable names for this example
train_features = X_train
train_labels = y_train
test_features = X_test
test_labels = y_test
print('Training Features Shape:', train_features.shape)
print('Training Labels Shape:', train_labels.shape)
print('Testing Features Shape:', test_features.shape)
print('Testing Labels Shape:', test_labels.shape)
```

```
Training Features Shape: (76595, 538)
Training Labels Shape: (76595,)
Testing Features Shape: (37726, 538)
Testing Labels Shape: (37726,)
```

```
In [69]: # Create Smaller Dataframe of 1000 Rows
train_features_1000 = pd.DataFrame.sample(X_train, n=1000, random_state=123)
train_labels_1000 = pd.DataFrame.sample(y_train,n=1000, random_state=123)
test_features_1000 = pd.DataFrame.sample(X_test,n=1000, random_state=123)
test_labels_1000 = pd.DataFrame.sample(y_test,n=1000, random_state=123)
print('Training Features 1000:', train_features_1000.shape)
print('Training Labels 1000:', train_labels_1000.shape)
print('Testing Features 1000:', test_features_1000.shape)
print('Testing Labels 1000:', test_labels_1000.shape)
```

```
Training Features 1000: (1000, 538)
Training Labels 1000: (1000,)
Testing Features 1000: (1000, 538)
Testing Labels 1000: (1000,)
```

```
In [70]: # Create Medium DataFrame of 5000 Rows
train_features_5000 = pd.DataFrame.sample(X_train, n=5000, random_state=123)
train_labels_5000 = pd.DataFrame.sample(y_train,n=5000, random_state=123)
test_features_5000 = pd.DataFrame.sample(X_test,n=5000, random_state=123)
test_labels_5000 = pd.DataFrame.sample(y_test,n=5000, random_state=123)
print('Training Features 5000:', train_features_5000.shape)
print('Training Labels 5000:', train_labels_5000.shape)
print('Testing Features 5000:', test_features_5000.shape)
print('Testing Labels 5000:', test_labels_5000.shape)
```

```
Training Features 5000: (5000, 538)
Training Labels 5000: (5000,)
Testing Features 5000: (5000, 538)
Testing Labels 5000: (5000,)
```

```
In [71]: # Original Code Commented - Using Pickled Models

# Use the random grid to search for best hyperparameters
# Define estimator - same parameters as base model

# rf = RandomForestClassifier(n_estimators=50, random_state=123 )
```

```
In [72]: # Create the FULL model based on the random_grid parameters
# rf_random = RandomizedSearchCV(estimator = rf, param_distributions = r
andom_grid, n_iter = 5, cv = 5, verbose=2, random_state=123, n_jobs = -
1)
```

```
In [73]: # Fit the SMALL random search model with 1000 rows
# rf_random_1000 = rf_random.fit(train_features_1000, train_labels_1000)
```

```
In [74]: # Fit the MEDIUM model with 5000 rows
# rf_random_5000 = rf_random.fit(train_features_5000, train_labels_5000)
```

```
In [75]: # Fit a model on all data using the same random_grid parameters
# rf_random_all = rf_random.fit(train_features, train_labels)
```

```
In [76]: #pickle.dump(rf_random_all,open("rf_random_all.pkl","wb"))
#pickle.dump(rf_random_1000,open("rf_random_1000.pkl","wb"))
#pickle.dump(rf_random_5000,open("rf_random_5000.pkl","wb"))
#pickle.dump(rf_base,open("rf_base.pkl","wb"))
#pickle.dump(rf_base_time,open("rf_base_time.pkl","wb"))
```

```
In [77]: # Load Pickled Models to avoid re-running models

rf_random_all = pickle.load( open("Pickle/rf_random_all.pkl", "rb"))
rf_random_1000 = pickle.load( open("Pickle/rf_random_1000.pkl", "rb"))
rf_random_5000 = pickle.load( open("Pickle/rf_random_5000.pkl", "rb"))
rf_base = pickle.load( open("Pickle/rf_base.pkl", "rb"))

# Load other stored variables
rf_base_time = pickle.load( open("Pickle/rf_base_time.pkl", "rb"))
```

```
In [78]: #
# Compare Random Grid Parameters among different datasets - 1000, 5000,
All
#####
####
```

```
In [79]: rf_random_1000.best_params_
```

```
Out[79]: {'n_estimators': 1000,
'min_samples_split': 10,
'min_samples_leaf': 2,
'max_features': 'auto',
'max_depth': None,
'bootstrap': False}
```

```
In [80]: rf_random_5000.best_params_
```

```
Out[80]: {'n_estimators': 1000,
'min_samples_split': 10,
'min_samples_leaf': 2,
'max_features': 'auto',
'max_depth': 1000,
'bootstrap': False}
```

```
In [81]: rf_random_all.best_params_
```

```
Out[81]: {'n_estimators': 1000,
'min_samples_split': 10,
'min_samples_leaf': 2,
'max_features': 'auto',
'max_depth': 1000,
'bootstrap': False}
```

```
In [82]: #  
# Determine Log-Loss and Accuracy for All RF Models  
#####
```

```
In [83]: # Generate Predictions for Base model  
rf_base_preds = rf_base.predict_proba(X_test)
```

```
In [84]: # Log-Loss and Accuracy for Base Model  
# import numpy as np  
# from sklearn.metrics import log_loss, accuracy_score  
rf_base_log_loss = log_loss(y_test, rf_base_preds[:,1]) # each column is  
# class probability,  
print(rf_base_log_loss)  
rf_base_accuracy = accuracy_score(y_test, np.rint(rf_base_preds[:,1]))  
print(rf_base_accuracy)  
  
0.24818377998398874  
0.9273975507607486
```

```
In [86]: # Generate Predictions for 1000 element tuned model  
rf_1000_preds = rf_random_1000.predict_proba(test_features_1000)
```

```
In [87]: # Log-Loss and Accuracy for 1000 Row Model  
rf_1000_log_loss = log_loss(test_labels_1000, rf_1000_preds[:,1]) # each  
# column is class probability,  
print(rf_1000_log_loss)  
rf_1000_accuracy = accuracy_score(test_labels_1000, np.rint(rf_1000_preds  
[:,1]))  
print(rf_1000_accuracy)  
  
0.48383002447131424  
0.786
```

```
In [88]: # Generate Predictions for 5000 element tuned model  
rf_5000_preds = rf_random_5000.predict_proba(test_features_5000)
```

```
In [89]: # Log-Loss and Accuracy for 5000 Row Model  
rf_5000_log_loss = log_loss(test_labels_5000, rf_5000_preds[:,1]) # each  
# column is class probability,  
print(rf_5000_log_loss)  
rf_5000_accuracy = accuracy_score(test_labels_5000, np.rint(rf_5000_preds  
[:,1]))  
print(rf_5000_accuracy)  
  
0.2587950210933267  
0.9172
```

```
In [90]: # Generate Predictions for Full Model  
rf_all_preds = rf_random_all.predict_proba(test_features)
```



```
In [91]: # Log-Loss and Accuracy for Full Model
rf_all_log_loss = log_loss(test_labels,rf_all_preds[:,1]) # each column
            is class probability,
print(rf_all_log_loss)
rf_all_accuracy = accuracy_score(test_labels,np.rint(rf_all_preds[:,1]))
print(rf_all_accuracy)
```

0.26495811383785384

0.9139585431797699

```
In [92]: #
# Feature Importance of Base RF Model
#####
```

```
In [93]: rf_base.feature_importances_
```

```
Out[93]: array([2.19576698e-02, 4.61408532e-03, 4.54027314e-03, 4.70482752e-03,
5.04434562e-03, 4.89174431e-03, 4.29335739e-03, 4.71446799e-03,
4.33285167e-03, 2.53118395e-02, 4.18768936e-03, 2.74401645e-02,
4.24549574e-03, 2.24728625e-02, 4.51155587e-03, 4.70138546e-03,
4.07981463e-03, 4.66154902e-03, 4.32649362e-03, 4.55287778e-03,
2.36252049e-02, 2.97495910e-03, 4.83270092e-03, 4.10554682e-03,
4.39553108e-03, 4.58738124e-03, 3.93261673e-03, 4.28856168e-03,
4.11493552e-03, 2.35251656e-02, 4.48150958e-03, 4.97092891e-03,
4.69097167e-03, 8.44852224e-04, 4.56074273e-03, 2.38307012e-02,
4.06558251e-03, 4.10675339e-03, 4.14507300e-03, 4.40307200e-03,
4.35607051e-03, 4.62930976e-03, 4.05495603e-03, 4.23632962e-03,
6.04834561e-02, 4.31834504e-03, 4.42748648e-03, 4.86906825e-03,
4.45527694e-03, 4.49465475e-03, 4.47505121e-03, 4.40967058e-03,
4.30586393e-03, 4.24420884e-03, 5.77902399e-03, 4.65919257e-03,
4.29646248e-03, 4.23596110e-03, 4.13769516e-03, 4.62747596e-03,
4.79847581e-03, 4.96326744e-03, 4.85468987e-03, 4.34816359e-03,
4.08267411e-03, 4.10669970e-03, 4.48577477e-03, 4.27239846e-03,
4.90572625e-03, 4.94663101e-03, 4.45542365e-03, 4.20347419e-03,
4.88923151e-03, 4.37896839e-03, 4.82268390e-03, 4.60666937e-03,
4.63525565e-03, 4.42760693e-03, 4.04569835e-03, 4.08931141e-03,
4.35296175e-03, 4.23361427e-03, 4.06210305e-03, 4.26940698e-03,
5.03295089e-03, 4.90955336e-03, 4.48772712e-03, 4.54793638e-03,
4.66487371e-03, 4.52246681e-03, 4.29532532e-03, 4.51421155e-03,
3.91441636e-03, 4.94528532e-03, 4.76524496e-03, 4.36378715e-03,
2.49368096e-02, 4.72404952e-03, 4.22808435e-03, 4.68432405e-03,
4.31596025e-03, 4.50018141e-03, 4.82522551e-03, 4.14196338e-03,
4.34990601e-03, 4.29001663e-03, 4.71185014e-03, 4.56621538e-03,
4.47170104e-03, 4.57825000e-03, 5.53908348e-03, 4.12908563e-03,
4.41852369e-03, 2.54004394e-05, 2.74933728e-05, 2.12694991e-05,
1.24306579e-04, 7.06309615e-05, 1.83493530e-04, 1.95827069e-04,
1.66573377e-04, 1.89412937e-04, 2.01096373e-04, 1.09923060e-04,
2.23111695e-04, 1.08342626e-04, 1.36453881e-04, 1.40313608e-04,
9.23674689e-04, 1.80819595e-04, 7.88653360e-05, 1.74448423e-04,
1.81885577e-04, 2.70236367e-04, 1.66474916e-04, 1.42727207e-04,
1.55435636e-04, 8.89766900e-05, 1.53067639e-04, 2.73087641e-04,
1.03984958e-04, 8.11105103e-05, 1.15089862e-04, 1.28179643e-04,
2.71396471e-04, 8.82736849e-05, 1.25203360e-04, 1.25073893e-04,
3.63750965e-04, 1.48598787e-04, 1.26851273e-04, 1.95659326e-04,
1.32535811e-04, 1.07272428e-04, 1.04785043e-04, 7.72072166e-05,
1.77091445e-04, 2.36757909e-04, 1.70837183e-04, 2.90568738e-04,
3.12715120e-04, 3.43341381e-04, 1.02091925e-04, 1.43551548e-04,
1.19334686e-04, 1.88758136e-04, 1.32031075e-04, 1.70739977e-04,
1.34041510e-04, 2.47020596e-04, 2.36518270e-04, 8.68803090e-05,
1.76367906e-04, 1.00568646e-04, 1.90118218e-04, 7.75785426e-04,
1.75561688e-04, 1.22049161e-03, 1.54836293e-03, 2.61477475e-03,
3.61264114e-03, 3.43330327e-03, 5.43763469e-04, 4.22599984e-05,
2.04852161e-03, 1.07197816e-03, 5.68018268e-04, 5.29238189e-04,
1.38754061e-03, 4.73326287e-03, 1.27595597e-03, 5.74684827e-04,
3.04144390e-05, 1.45258311e-05, 2.68055121e-03, 2.33255330e-04,
1.07274183e-03, 6.86160049e-04, 7.46941091e-04, 0.00000000e+00,
2.68969182e-03, 7.07513138e-04, 2.27735842e-03, 2.24227069e-03,
2.46207353e-03, 2.32227765e-03, 2.33086267e-03, 2.52010851e-03,
2.37228516e-03, 2.29398076e-03, 2.48736925e-03, 2.66476152e-03,
2.13190093e-03, 2.28080932e-03, 1.20596312e-05, 3.13181928e-05,
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 4.53716157e-04,
4.62617392e-04, 3.22426617e-05, 1.32723410e-04, 0.00000000e+00,
3.96113319e-05, 7.24101742e-05, 4.80447222e-07, 1.59100020e-04,
```

1.00228622e-04, 0.00000000e+00, 7.62196353e-05, 9.13298174e-04,  
2.29698122e-06, 1.56694305e-05, 0.00000000e+00, 9.42177642e-04,  
0.00000000e+00, 2.51536688e-05, 9.65568867e-05, 7.24120813e-07,  
1.75796957e-04, 0.00000000e+00, 1.02147554e-05, 0.00000000e+00,  
2.31585403e-06, 5.78321660e-05, 2.23130924e-05, 6.82549862e-06,  
8.58384446e-04, 5.74466965e-04, 5.14431604e-04, 1.02733948e-04,  
1.31125866e-05, 0.00000000e+00, 0.00000000e+00, 5.13112954e-05,  
1.21673683e-06, 1.20252649e-05, 0.00000000e+00, 2.50264158e-05,  
4.63700364e-04, 2.60120604e-03, 3.11364563e-04, 4.93507569e-05,  
1.23851821e-03, 3.46982077e-05, 4.79993237e-05, 6.57883778e-06,  
9.26484896e-05, 0.00000000e+00, 0.00000000e+00, 3.31148179e-05,  
0.00000000e+00, 5.61997005e-05, 1.10851826e-04, 0.00000000e+00,  
0.00000000e+00, 1.79237738e-05, 1.86935927e-04, 1.19757564e-03,  
1.49372707e-05, 1.91049485e-04, 0.00000000e+00, 1.19433666e-04,  
8.53069442e-07, 0.00000000e+00, 9.14139534e-06, 7.14993594e-07,  
1.79725404e-03, 0.00000000e+00, 0.00000000e+00, 1.87453695e-04,  
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 7.23094797e-06,  
6.76028780e-04, 1.86917371e-05, 2.88857284e-04, 1.35746406e-03,  
5.10741293e-04, 3.77701227e-05, 4.21926473e-05, 1.93095000e-05,  
5.43978123e-05, 1.39638366e-03, 1.13187976e-03, 1.69381039e-05,  
5.39827564e-04, 4.40109539e-04, 0.00000000e+00, 4.42090457e-05,  
2.72504858e-05, 1.72614436e-05, 8.07719691e-04, 5.93405813e-04,  
3.00543378e-05, 0.00000000e+00, 1.72129229e-06, 4.38762351e-04,  
0.00000000e+00, 0.00000000e+00, 0.00000000e+00, 0.00000000e+00,  
7.50338005e-04, 0.00000000e+00, 2.74232348e-03, 0.00000000e+00,  
2.55836996e-04, 0.00000000e+00, 3.72568845e-04, 2.21190838e-04,  
1.43084138e-07, 0.00000000e+00, 2.03106847e-04, 4.14633004e-05,  
4.47640158e-03, 5.46458158e-03, 6.95946871e-03, 0.00000000e+00,  
2.93317877e-03, 1.73990829e-03, 0.00000000e+00, 2.99929110e-03,  
2.56776743e-06, 1.70231687e-05, 0.00000000e+00, 0.00000000e+00,  
7.52471660e-06, 2.43865672e-04, 2.66432655e-04, 1.72003022e-05,  
2.76946338e-03, 3.87356054e-05, 2.86155345e-03, 8.84101574e-05,  
4.00973227e-03, 2.34966438e-03, 9.46756254e-04, 1.32674100e-03,  
7.47842613e-05, 1.50761212e-05, 5.60172094e-04, 8.50274306e-04,  
1.14709526e-04, 6.07251902e-04, 0.00000000e+00, 8.09251666e-04,  
1.83656204e-05, 3.58782012e-04, 2.03227006e-04, 5.64559573e-04,  
1.28860827e-05, 3.10084072e-03, 2.84709424e-03, 2.96224209e-03,  
9.68176636e-05, 7.26218018e-04, 2.19342257e-03, 3.15095697e-03,  
2.19159580e-03, 2.84088295e-03, 2.88437032e-03, 2.96774490e-03,  
3.09536333e-03, 7.05958777e-04, 1.15147677e-04, 4.08359096e-03,  
3.22983596e-03, 2.58329173e-04, 1.88723942e-03, 9.34313633e-04,  
7.57055648e-04, 1.76222247e-03, 1.49479948e-03, 3.08634085e-03,  
7.13091307e-04, 1.44654322e-03, 2.20296468e-03, 1.39830127e-03,  
1.12689287e-03, 1.23370596e-03, 5.72953441e-04, 1.88697845e-03,  
1.16985829e-03, 1.46906566e-03, 8.77790809e-04, 1.27767507e-03,  
4.11196560e-04, 1.16643982e-03, 1.33464864e-03, 7.74161609e-04,  
2.88282305e-04, 1.56963778e-08, 6.58285595e-04, 1.66247279e-03,  
2.01795038e-04, 5.57773878e-04, 1.19065500e-03, 8.59052095e-04,  
5.14626666e-04, 3.26117624e-04, 5.58825139e-04, 2.21572557e-07,  
5.84951695e-04, 3.85437459e-04, 1.76457117e-04, 2.17623585e-04,  
9.21530546e-05, 3.21602650e-03, 1.22140220e-05, 9.60876638e-04,  
1.60124049e-04, 3.09974776e-04, 9.75534992e-04, 2.01688740e-04,  
8.31186668e-05, 9.71383991e-04, 3.57265403e-04, 1.25548067e-04,  
5.10551899e-04, 4.94192270e-04, 3.61042372e-04, 6.57704490e-04,  
7.08350161e-04, 6.41547994e-04, 7.26995550e-04, 2.72957610e-04,  
8.33271017e-04, 3.85393631e-04, 9.61226200e-05, 1.01715035e-03,  
3.61970603e-04, 4.12161562e-04, 7.12664832e-04, 4.26350366e-04,

```

3.09721409e-04, 6.00479217e-04, 8.51278834e-05, 1.36145431e-03,
5.99698843e-04, 3.72760455e-04, 6.86691465e-04, 3.75835536e-04,
1.36797109e-03, 2.47316479e-04, 1.02583890e-03, 3.23943725e-04,
4.87512832e-04, 5.91838392e-04, 3.65676160e-04, 3.69384385e-04,
4.17627861e-05, 5.08768857e-04, 7.78862369e-04, 8.68579165e-04,
2.50864184e-04, 7.93733299e-05, 2.64925241e-04, 9.41540318e-04,
4.17095219e-04, 4.02885951e-04, 1.60405536e-04, 4.96035495e-04,
4.28745449e-04, 1.42545366e-03, 7.53628543e-04, 6.26934009e-04,
1.77717844e-03, 4.61356253e-04, 5.11173338e-04, 4.01239311e-04,
3.67931304e-04, 4.21301233e-04, 2.35173402e-04, 2.49097441e-04,
5.74008478e-04, 4.82927529e-04, 1.04483759e-03, 7.18448703e-04,
1.37272816e-03, 1.82682680e-06, 3.31779683e-04, 6.97077519e-04,
4.24798918e-04, 4.37114327e-04, 8.57595020e-04, 4.95756146e-04,
5.86053205e-04, 1.21320890e-03, 4.93604290e-04, 2.52189488e-04,
7.74686488e-04, 3.62113573e-04, 2.91734814e-04, 5.98962065e-04,
1.01921155e-03, 2.43672730e-04, 9.76607112e-04, 1.24287848e-03,
2.86509193e-04, 5.14574235e-04, 1.17435983e-03, 1.18877379e-03,
4.67374030e-04, 6.28369628e-04, 2.37936407e-04, 7.97930511e-04,
6.62439286e-04, 5.69390890e-04, 4.97797317e-04, 4.72505073e-04,
5.31497414e-04, 1.00473644e-03, 3.97094551e-04, 4.76147072e-04,
5.54053070e-04, 5.95136452e-04])

```

```

In [94]: # Top 10 Features for Base Model
importances = rf_base.feature_importances_
std = np.std([tree.feature_importances_ for tree in rf_base.estimators_],
              axis=0)
indices = np.argsort(importances)[::-1]

```

```

In [95]: #import numpy as np
import matplotlib.pyplot as plt

# Print the feature ranking
# print("Feature ranking:")

#for f in range(X_train.shape[1]):
#    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

print("Top Ten Features:")

for f in range(0, 10):
    print("%d. feature %d (%f)" % (f + 1, indices[f], importances[indices[f]]))

```

```

Top Ten Features:
1. feature 44 (0.060483)
2. feature 11 (0.027440)
3. feature 9 (0.025312)
4. feature 96 (0.024937)
5. feature 35 (0.023831)
6. feature 20 (0.023625)
7. feature 29 (0.023525)
8. feature 13 (0.022473)
9. feature 0 (0.021958)
10. feature 338 (0.006959)

```

```
In [96]: #####
# RESULTS
#####
```

```
In [97]: # Build Summary Table for All Model Types
data = [['XGBoost', 'Base', xgb_base_log_loss, xgb_base_accuracy, xgb_base_time],
        ['Random Forest Complete Dataset', 'Base', rf_base_log_loss, rf_base_accuracy, rf_base_time],
        ['Random Forest Complete Dataset', 'Tuned', rf_all_log_loss, rf_all_accuracy, 90],
        ['Random Forest 1000 Entries', 'Tuned', rf_1000_log_loss, rf_1000_accuracy, 2],
        ['Random Forest 5000 Entries', 'Tuned', rf_5000_log_loss, rf_5000_accuracy, 7]]
```

```
In [98]: #from tabulate import tabulate
print (tabulate(data, headers=["Model", "Tuning", "Log-Loss", "Accuracy", "Wall Time"]))
```

Model	Tuning	Log-Loss	Accuracy	Wall Time
XGBoost	Base	0.585187	0.766395	7.89
Random Forest Complete Dataset	Base	0.248184	0.927398	41.22
Random Forest Complete Dataset	Tuned	0.264958	0.913959	90
Random Forest 1000 Entries	Tuned	0.48383	0.786	2
Random Forest 5000 Entries	Tuned	0.258795	0.9172	7

```
In [ ]: ##
# Visualizations
#####
```

```
In [ ]: # Single Tree Visualization
# All Credit to https://www.kaggle.com/willkoehrsen/a-complete-introduction-and-walkthrough#Visualize-Single-Decision-Tree

#model = RandomForestClassifier(max_depth = 3, n_estimators=10)
#model.fit(train_selected, train_labels)
model = rf_base
estimator_limited = model.estimators_[1]
estimator_limited
```

```
In [ ]: train_selected = X_train
```

```
In [ ]: #from sklearn.tree import export_graphviz

export_graphviz(estimator_limited, out_file='Images/1_tree.dot', feature
_names = train_selected.columns,
                rounded = True, proportion = False, precision = 2, fille
d = True)
```

```
In [ ]: # Convert .dot file to .png - commented out because run time is long.
# Using .png in write-up
#import os
#os.environ["PATH"] += os.pathsep + 'C:/Program Files/Graphviz/bin/'
#os.system('dot -Tpng tree_limited.dot -o tree_limited.png')
```

```
In [99]: ### Create Small Tree with 3 Levels to be able to visualize
# Look at parameters used by our current forest
print('Parameters currently in use:\n')
pprint(rf_base.get_params())
```

Parameters currently in use:

```
{'bootstrap': True,
 'ccp_alpha': 0.0,
 'class_weight': None,
 'criterion': 'gini',
 'max_depth': None,
 'max_features': 'auto',
 'max_leaf_nodes': None,
 'max_samples': None,
 'min_impurity_decrease': 0.0,
 'min_impurity_split': None,
 'min_samples_leaf': 1,
 'min_samples_split': 2,
 'min_weight_fraction_leaf': 0.0,
 'n_estimators': 50,
 'n_jobs': None,
 'oob_score': False,
 'random_state': 123,
 'verbose': 0,
 'warm_start': False}
```

```
In [ ]: # Limit depth of tree to 3 levels to be able to see details.
rf_small = RandomForestClassifier(n_estimators=10, max_depth = 3)
rf_small.fit(train_features, train_labels)
# Extract the small tree
tree_small = rf_small.estimators_[5]

# Save the tree as a png image - Commented out to save time since already done
# export_graphviz(tree_small, out_file = 'Images\\small_tree.dot', feature_names = train_selected.columns, rounded = True, precision = 1)
# (graph, ) = pydot.graph_from_dot_file('small_tree.dot')
# graph.write_png('small_tree.png');
```

```
In [ ]:
```

In [ ]:

In [ ]:

In [ ]:

In [ ]: ##### Old Code Below Here #####

```
In [37]: #from sklearn.ensemble import RandomForestClassifier
rf = RandomForestClassifier(n_estimators=50)

start = time.time()
rf.fit(X_train, y_train)

end = time.time()
rf_time=round((end-start),2)
rf_time
```

Out[37]: 31.57

```
In [40]: preds = rf.predict_proba(X_test)
```

```
In [41]: rf_base_log_loss = log_loss(y_test,preds[:,1]) # each column is class pr
obability,
print(rf_base_log_loss)
rf_base_accuracy = accuracy_score(y_test,np rint(preds[:,1]))
print(rf_base_accuracy)

0.4925212965112307
0.7761755818268569
```

```
In [43]: data = [['XGBoost',xgb_base_log_loss, xgb_base_accuracy, xgb_time],
                ['LinearSVC', "N/A", svm_base_accuracy, LinearSVC_time],
                ['Random Forest',rf_base_log_loss, rf_base_accuracy, rf_time]]
```

```
In [44]: print(data)
```

```
[['XGBoost', 0.5851586336031108, 0.7655198006679743, 10.97], ['LinearSV
C', 'N/A', 0.528547951015215, 61.67], ['Random Forest', 0.4925212965112
307, 0.7761755818268569, 31.57]]
```

```
In [45]: rowlen=len(data)
print(rowlen)
```

3



```
In [46]: import numpy as np
rownums=np.arange(0,rowlen,1)
rownums=rownums+1

headers=['Model', 'Log-Loss', 'Accuracy', 'Wall Time']
print(headers)
print(rownums)

['Model', 'Log-Loss', 'Accuracy', 'Wall Time']
[1 2 3]
```

```
In [47]: print(pd.DataFrame(data, rownums,headers))
```

	Model	Log-Loss	Accuracy	Wall Time
1	XGBoost	0.585159	0.765520	10.97
2	LinearSVC	N/A	0.528548	61.67
3	Random Forest	0.492521	0.776176	31.57

```
In [48]: # Output table of the models results


print (tabulate(data, headers=["Model", "Log-Loss", "Accuracy", "Wall Time"] ))
```

Model	Log-Loss	Accuracy	Wall Time
XGBoost	0.5851586336031108	0.76552	10.97
LinearSVC	N/A	0.528548	61.67
Random Forest	0.4925212965112307	0.776176	31.57

```
In [ ]: #####
#. Extra code for loading and plotting
#####
```

```
In [42]: # Plot the Random Forest model vs SVM
ax = plt.gca()
rfc_disp = plot_roc_curve(rf, X_test, y_test, ax=ax, alpha=0.8)
svm_disp.plot(ax=ax, alpha=0.8)
```

Out[42]: <sklearn.metrics.\_plot.roc\_curve.RocCurveDisplay at 0x1a23145f50>



```
In [31]: svc = LinearSVC(random_state=123)
svc.fit(X_train, y_train)
svc_disp = plot_roc_curve(svc, X_test, y_test)
```

```
-----
-----
NameError                                Traceback (most recent call 1
ast)
<ipython-input-31-863beec980c9> in <module>
----> 1 svc = SVCLinear(random_state=42)
      2 svc.fit(X_train, y_train)
      3 svc_disp = plot_roc_curve(svc, X_test, y_test)

NameError: name 'SVCLinear' is not defined
```

```
In [ ]: # Random forest classifier
rfc = RandomForestClassifier(random_state=42)
rfc.fit(X_train, y_train)

ax = plt.gca()
rfc_disp = plot_roc_curve(rfc, X_test, y_test, ax=ax, alpha=0.8)
svm_disp.plot(ax=ax, alpha=0.8)
```

```
In [ ]: X, y = datasets.make_classification(random_state=0)
        clf = svm.SVC(random_state=0)
        clf.fit(X_train, y_train)
        SVC(random_state=0)
        metrics.det_curve(clf, X_test, y_test)
        plt.show()
```

```
In [ ]: import scikitplot as skplt
        rf = RandomForestClassifier()
        lr = LogisticRegression()
        nb = GaussianNB()
        svm = LinearSVC()
        rf_probas = rf.fit(X_train, y_train).predict_proba(X_test)
        lr_probas = lr.fit(X_train, y_train).predict_proba(X_test)
        nb_probas = nb.fit(X_train, y_train).predict_proba(X_test)
        svm_scores = svm.fit(X_train, y_train).decision_function(X_test)
        probas_list = [rf_probas, lr_probas, nb_probas, svm_scores]
        clf_names = ['Random Forest', 'Logistic Regression',
                     'Gaussian Naive Bayes', 'Support Vector Machine']
        skplt.metrics.plot_calibration_curve(y_test,
                                           probas_list,
                                           clf_names)
        plt.show()
```

```
In [ ]: from sklearn import metrics
        metrics.det_curve(clf, X_test, y_test)
        plt.show()
```

```
In [ ]:
```